TCP Overview

TCP/IP Protocol Stack

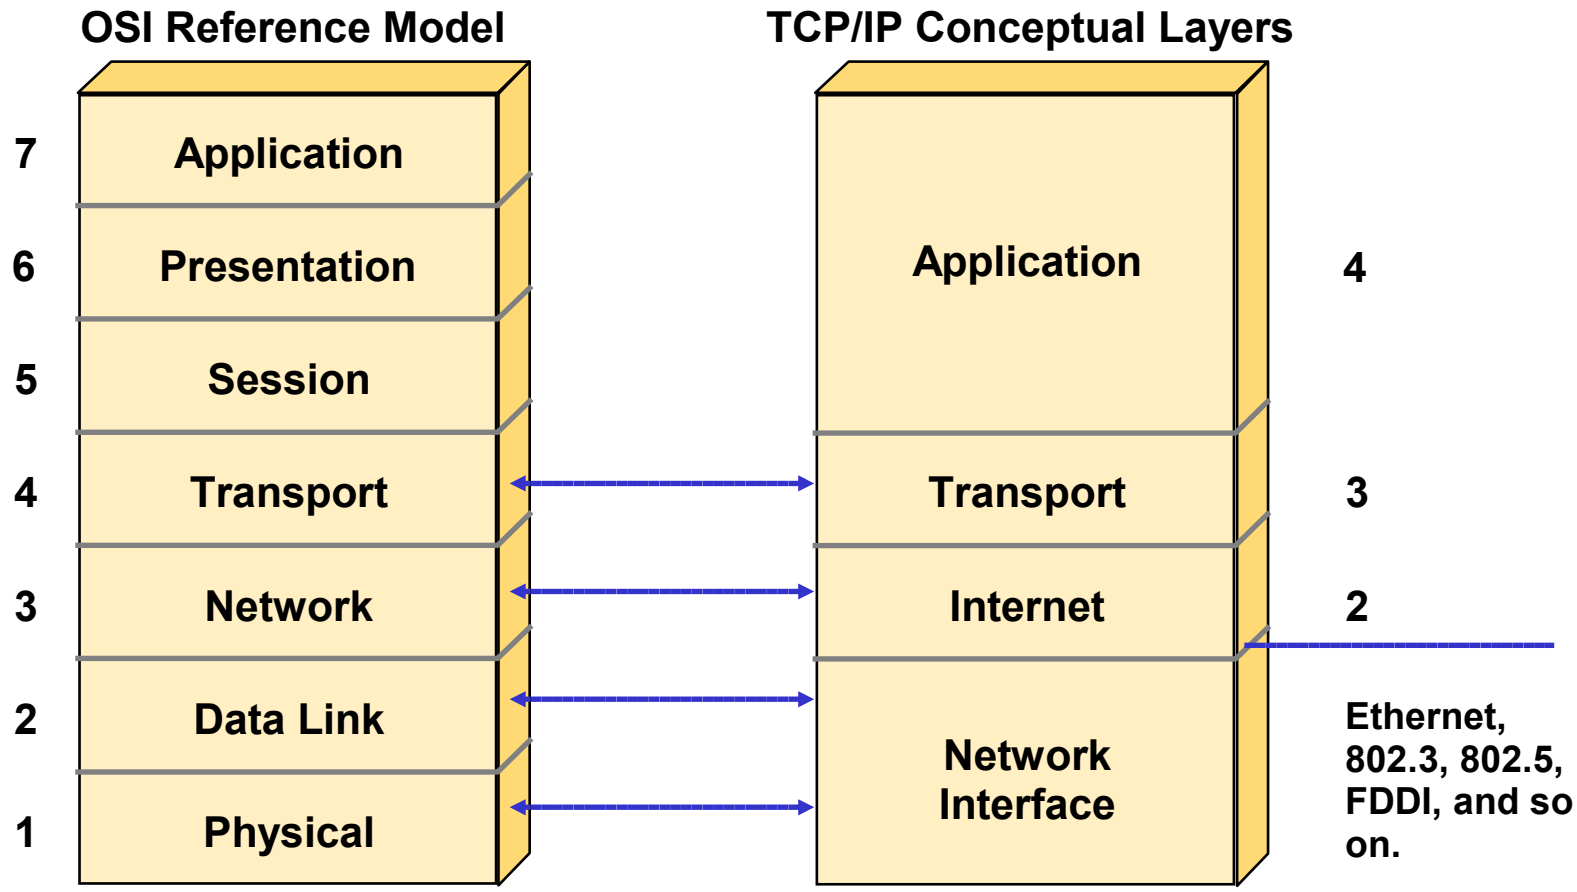

Transport Layer Overview

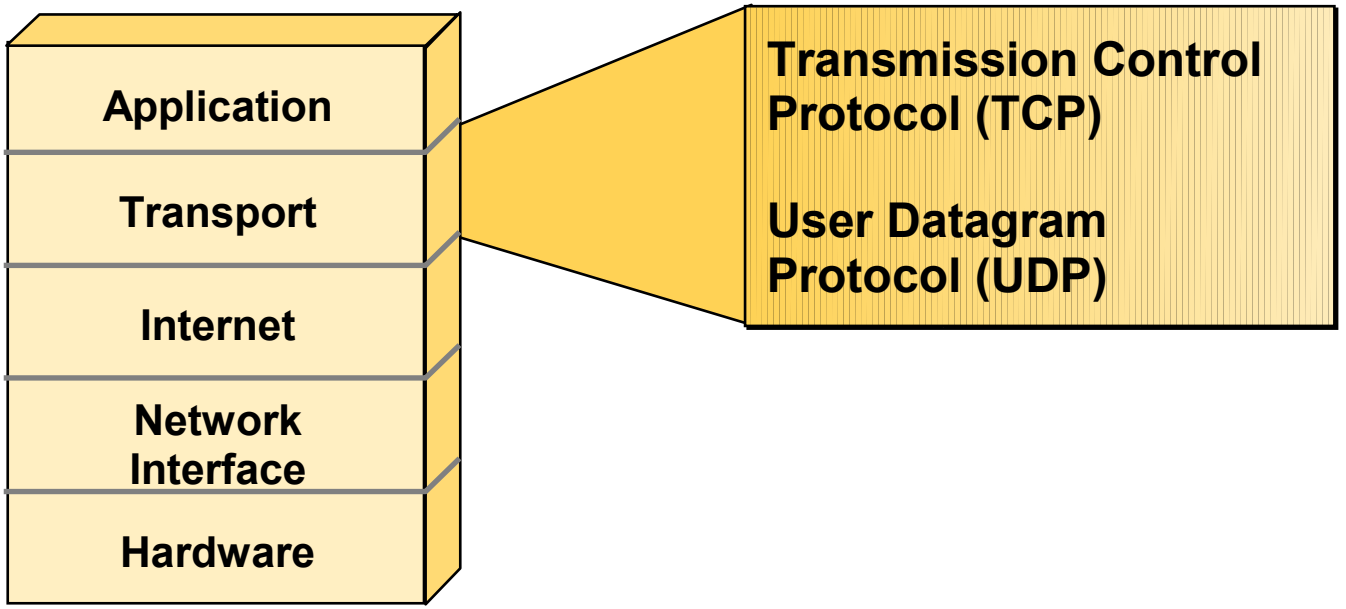

TCP

- Defined in RFC 793
- Reliable
  - Acknowledgments
  - Guarantee of packet delivery
  - Delayed Ack – Piggybacking
  - Reassembly of out of order data
  - Discards duplicates caused by IP
  - Provided end-to-end flow control
  - finite buffer size

TCP

- Connection oriented
  - Segments are dependent
  - Maintain state information of segments
  - Segments can take different routes
  - Segments are delivered in order to the application layer
- Full Duplex

TCP Segment Format

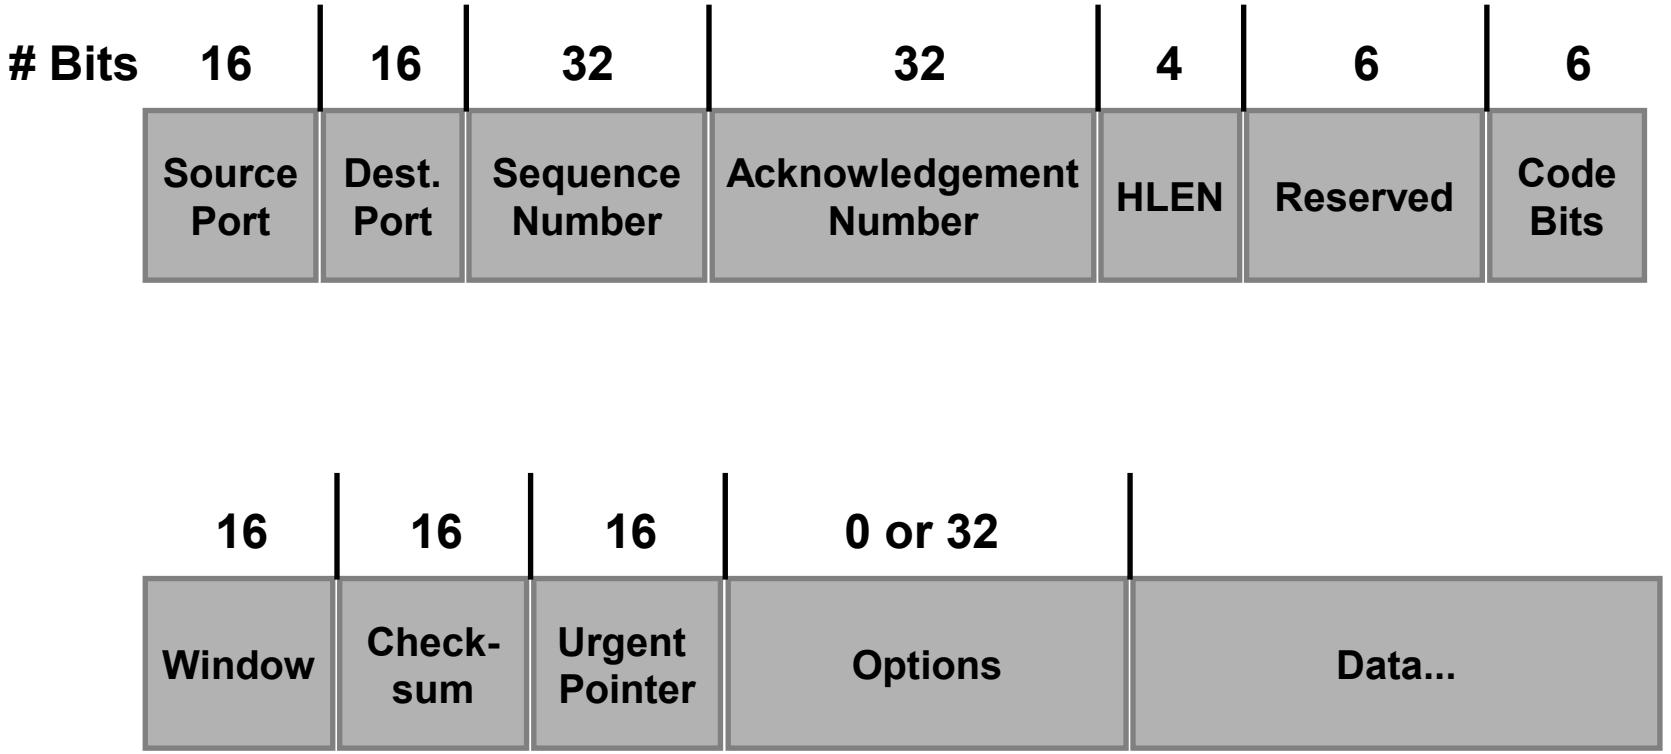

Port Numbers

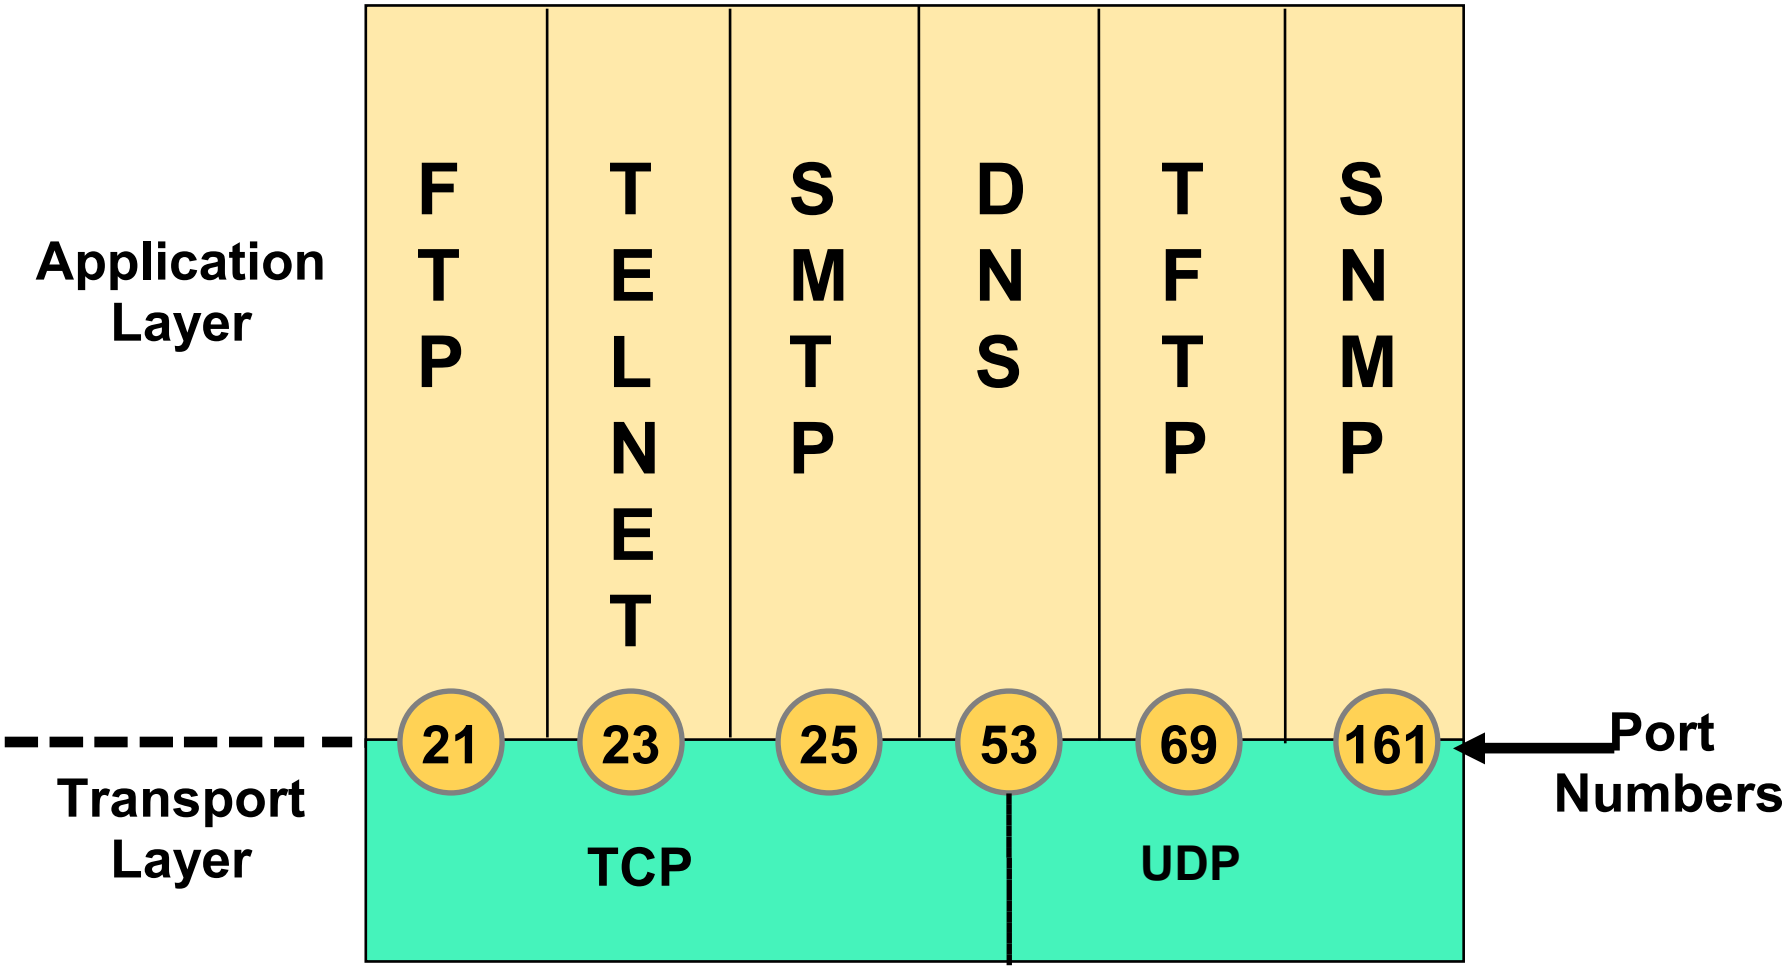

TCP Port Numbers

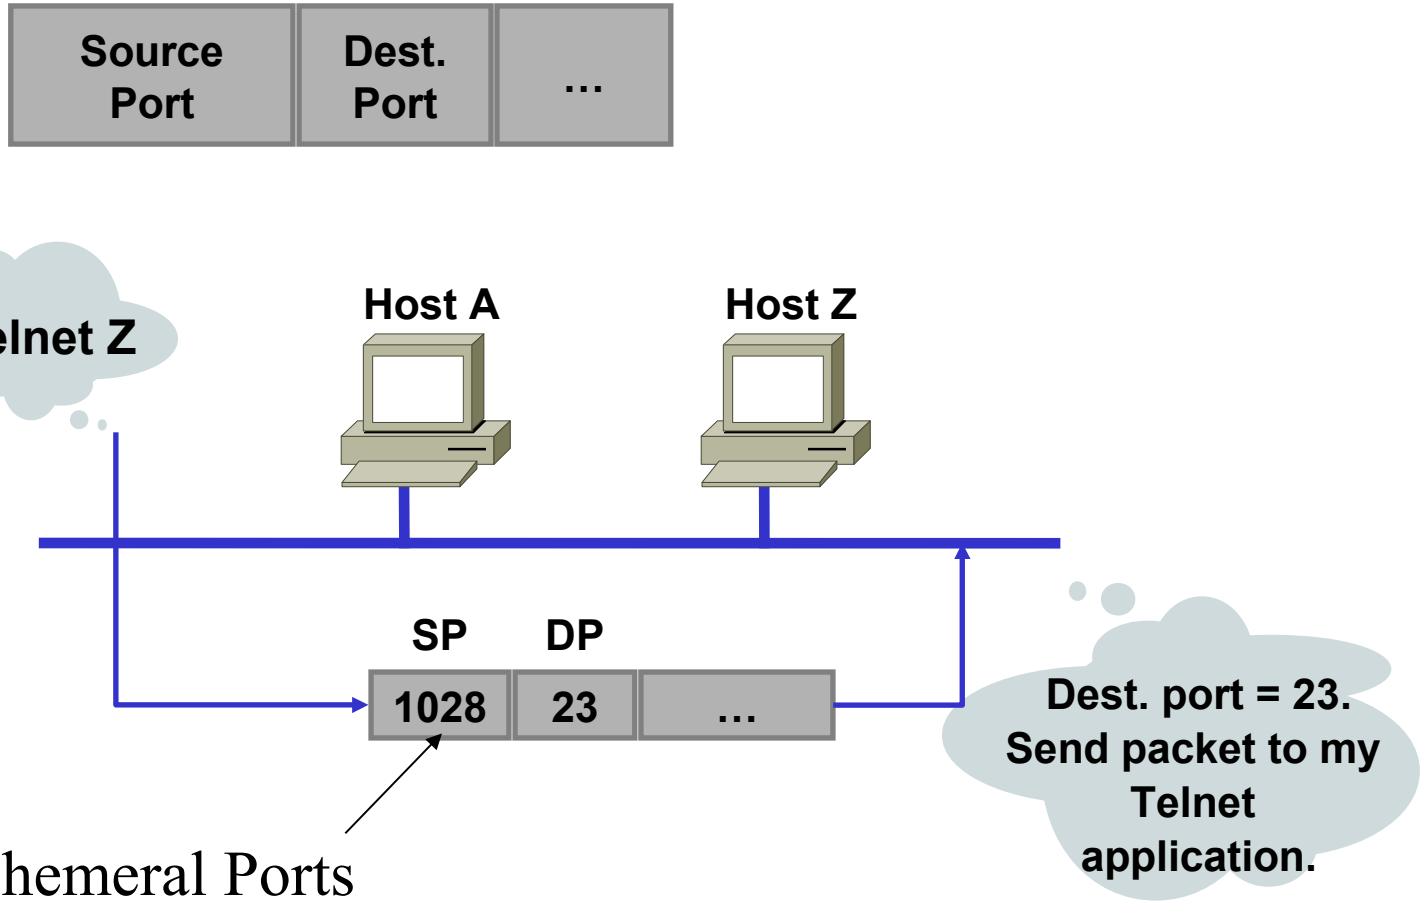

TCP Handshake/Open Connection

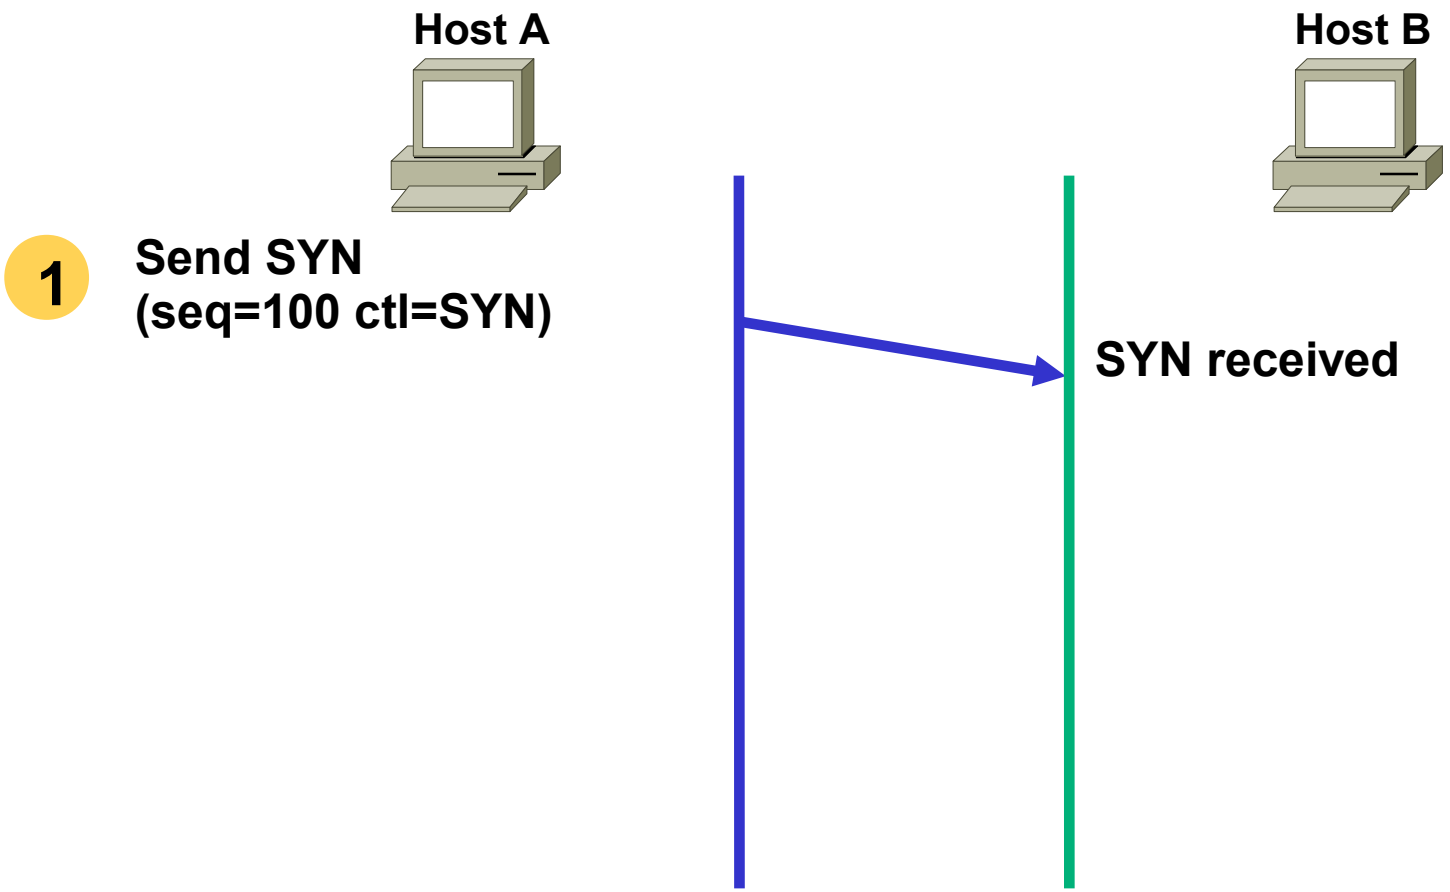

TCP Handshake/Open Connection

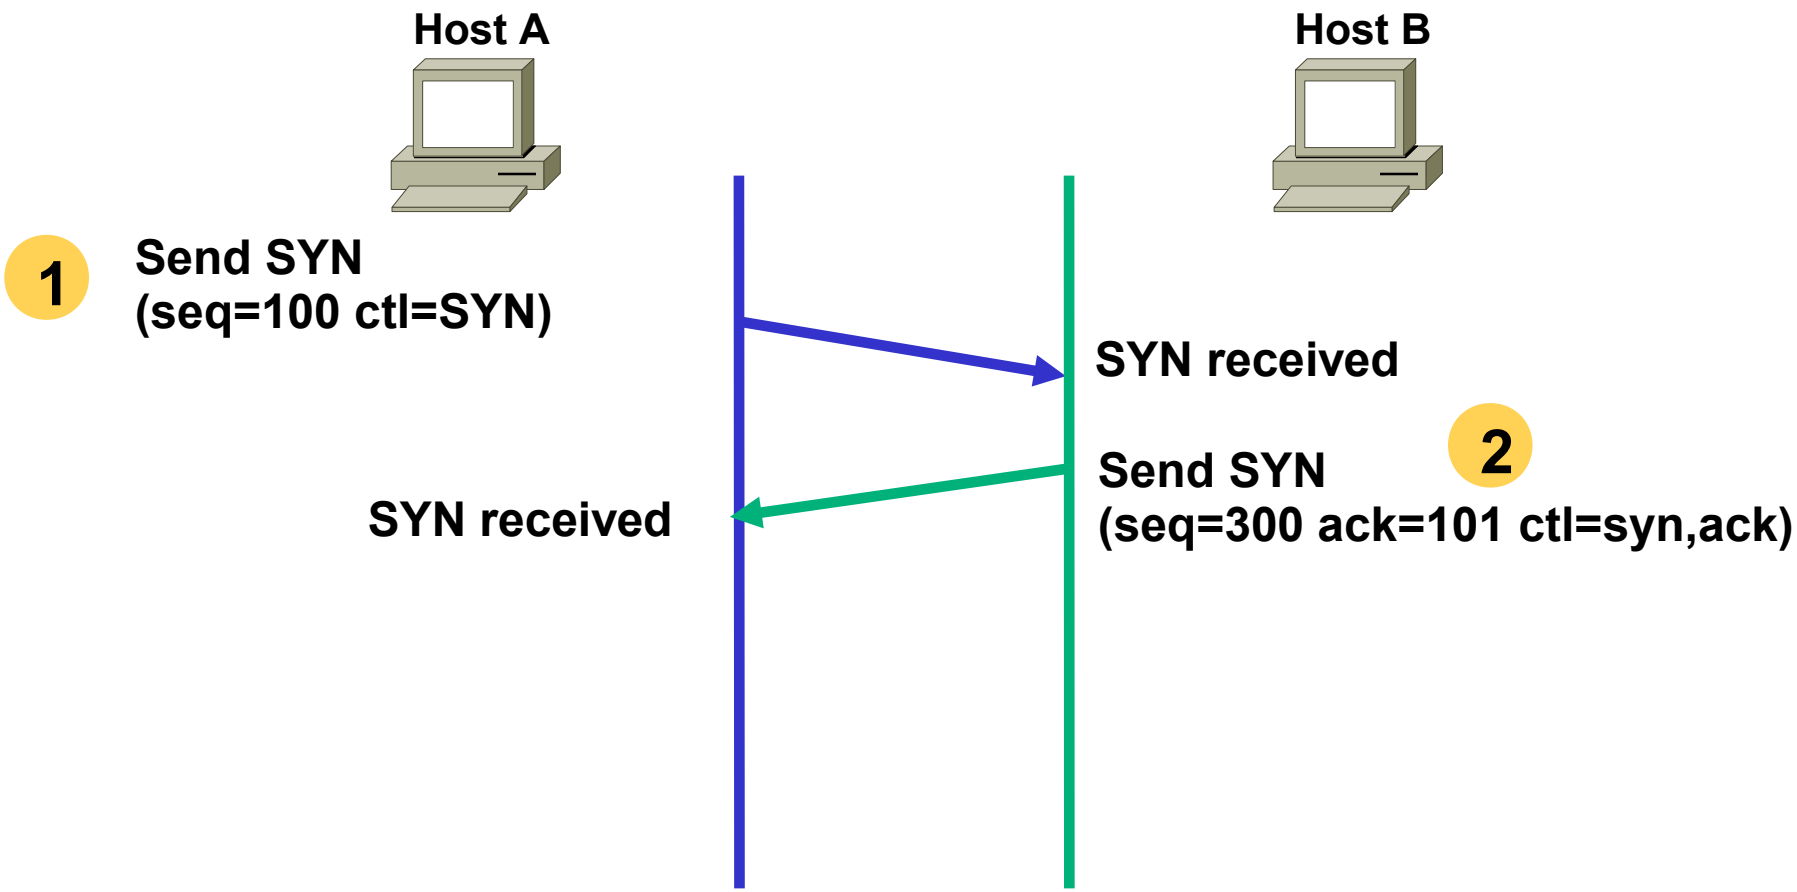

TCP Handshake/Open Connection

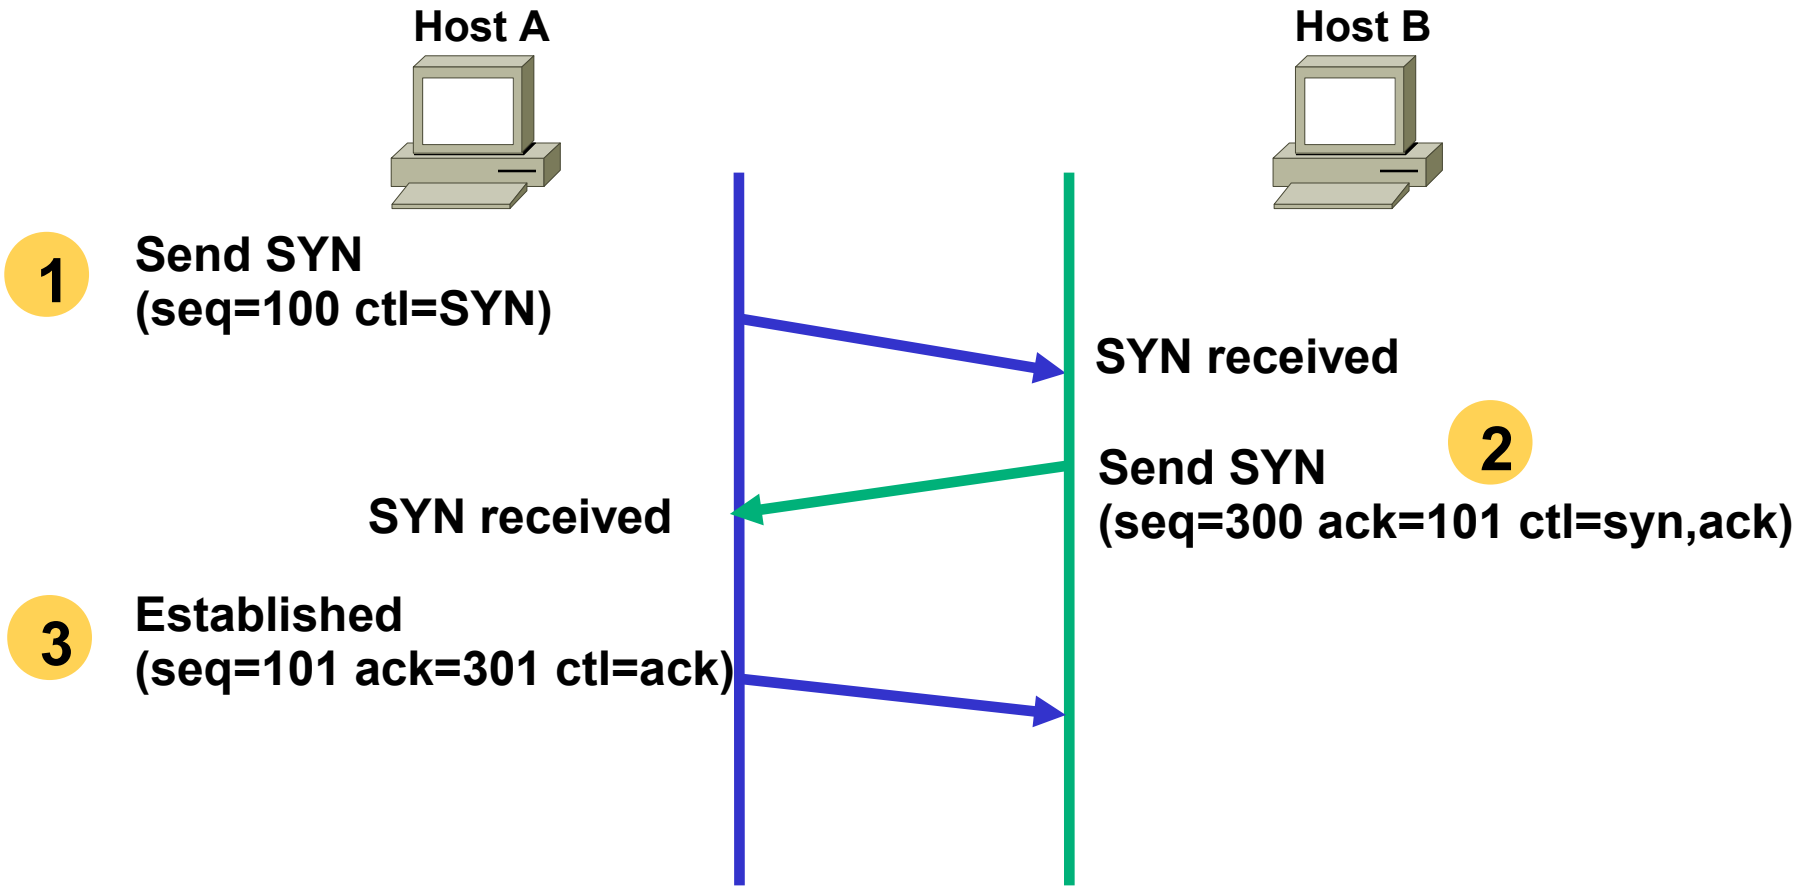

TCP Handshake/Open Connection

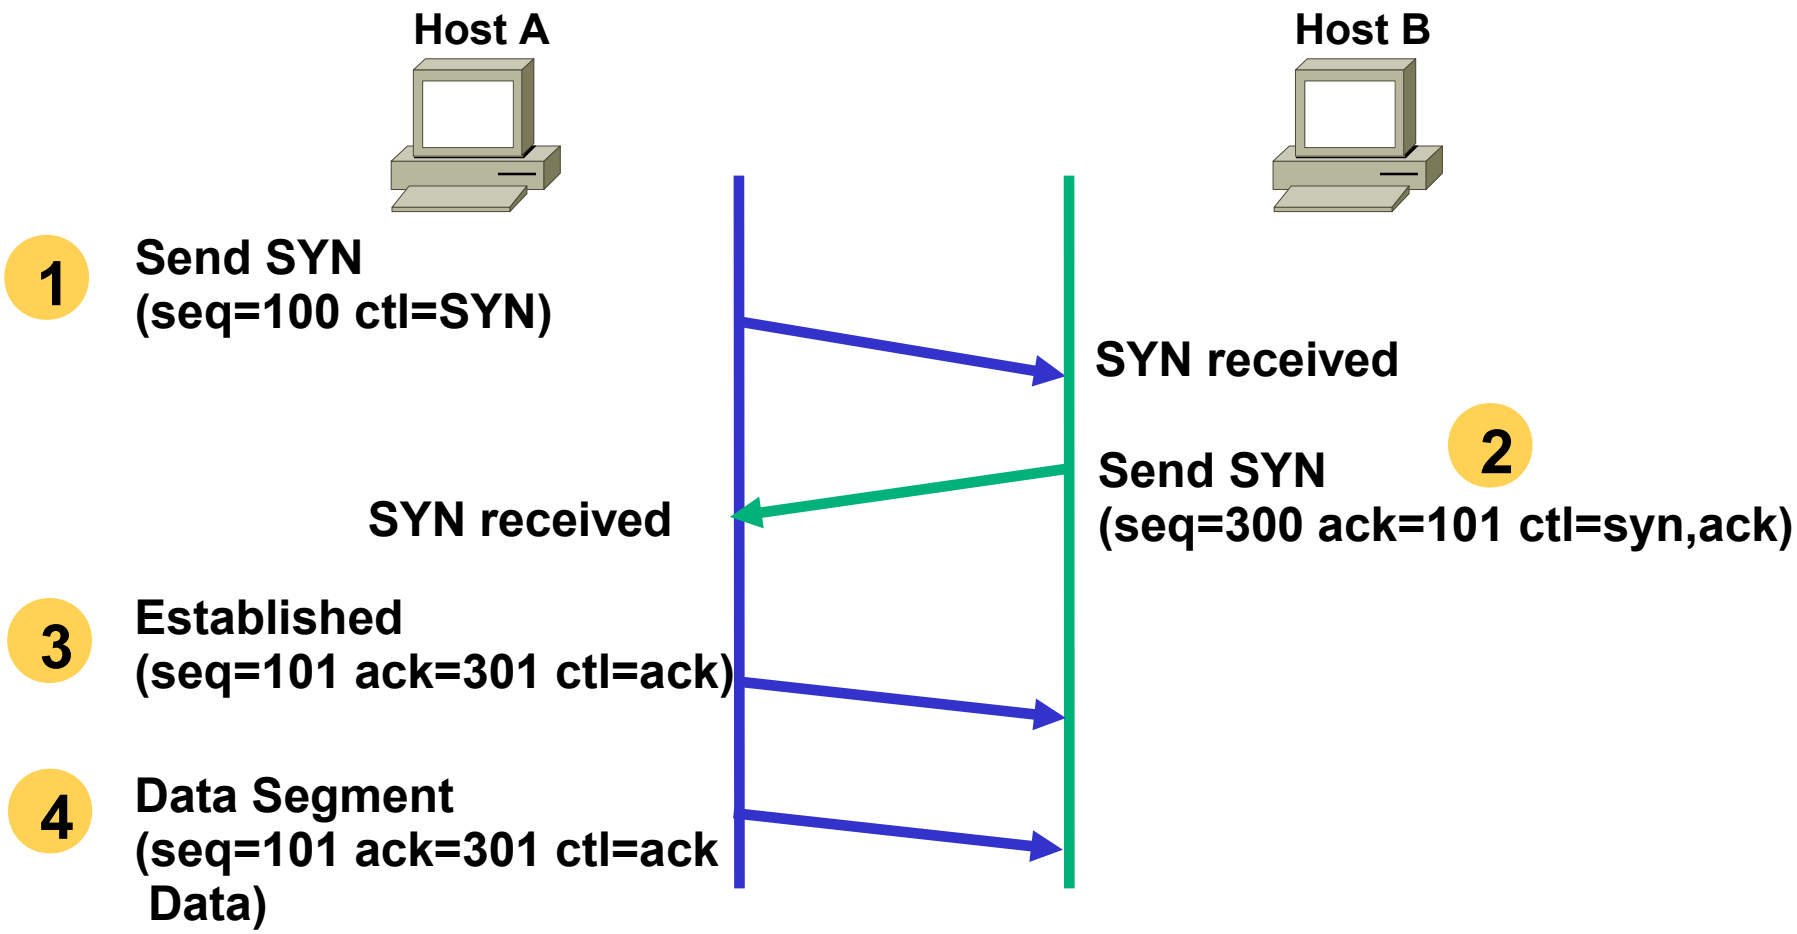

TCP Simple Acknowledgment

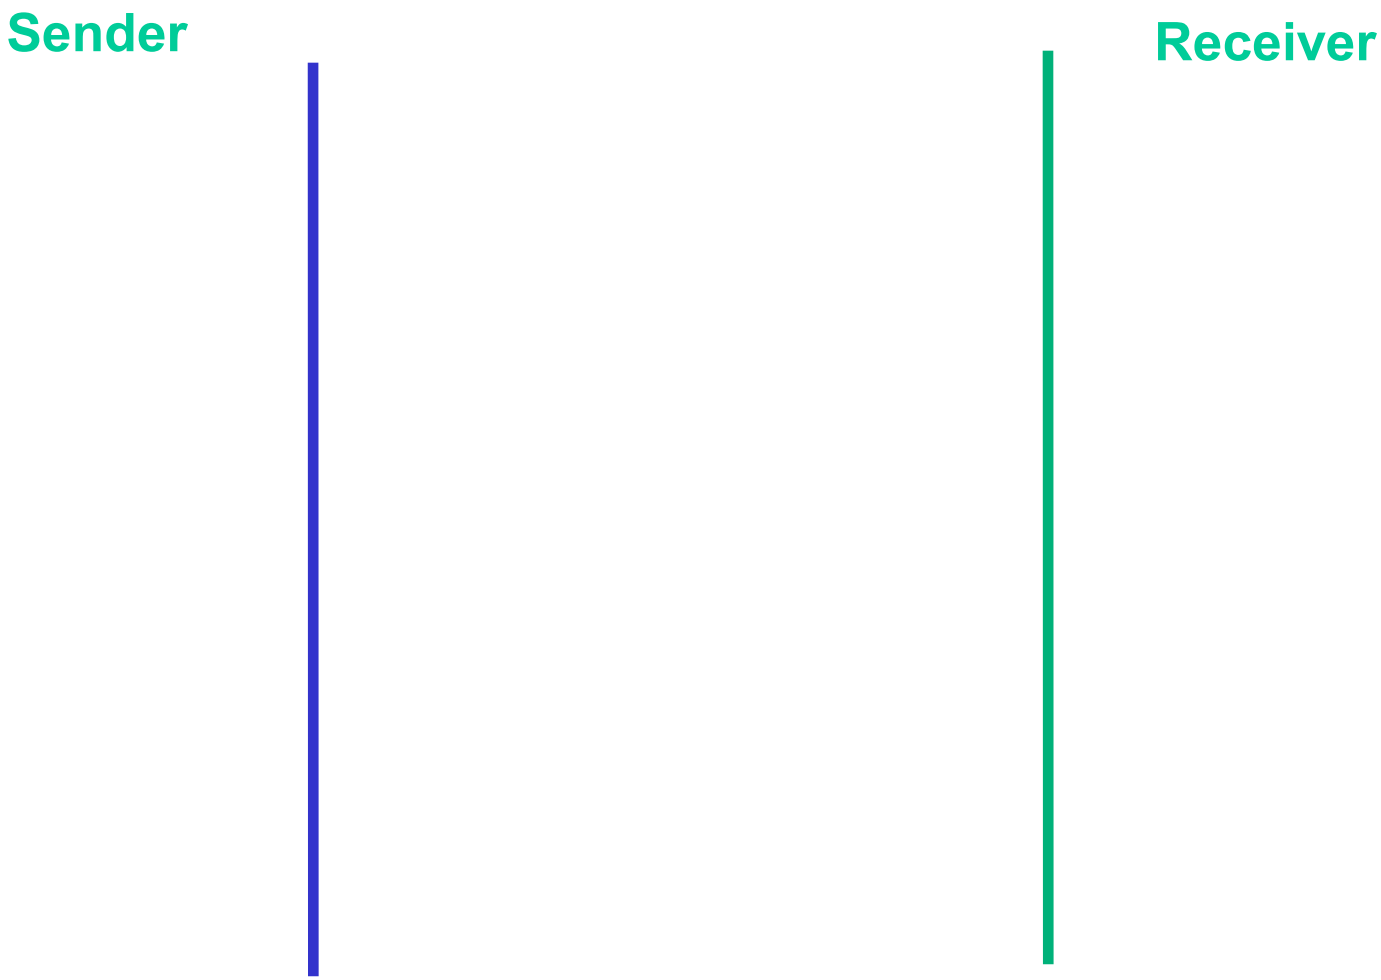

- Window size = 1

TCP Simple Acknowledgment

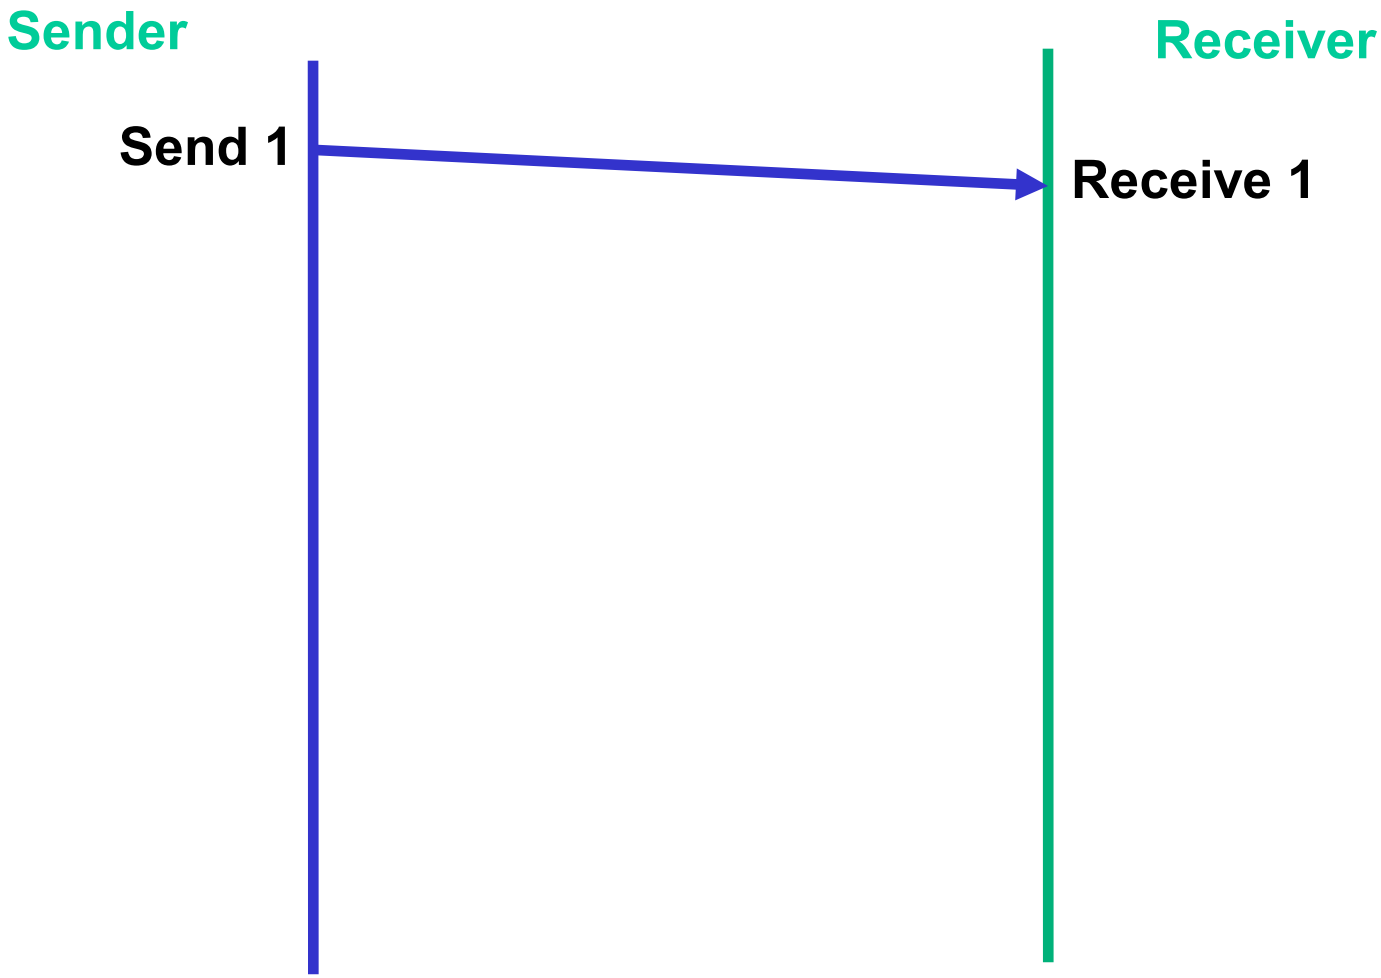

- Window size = 1

TCP Simple Acknowledgment

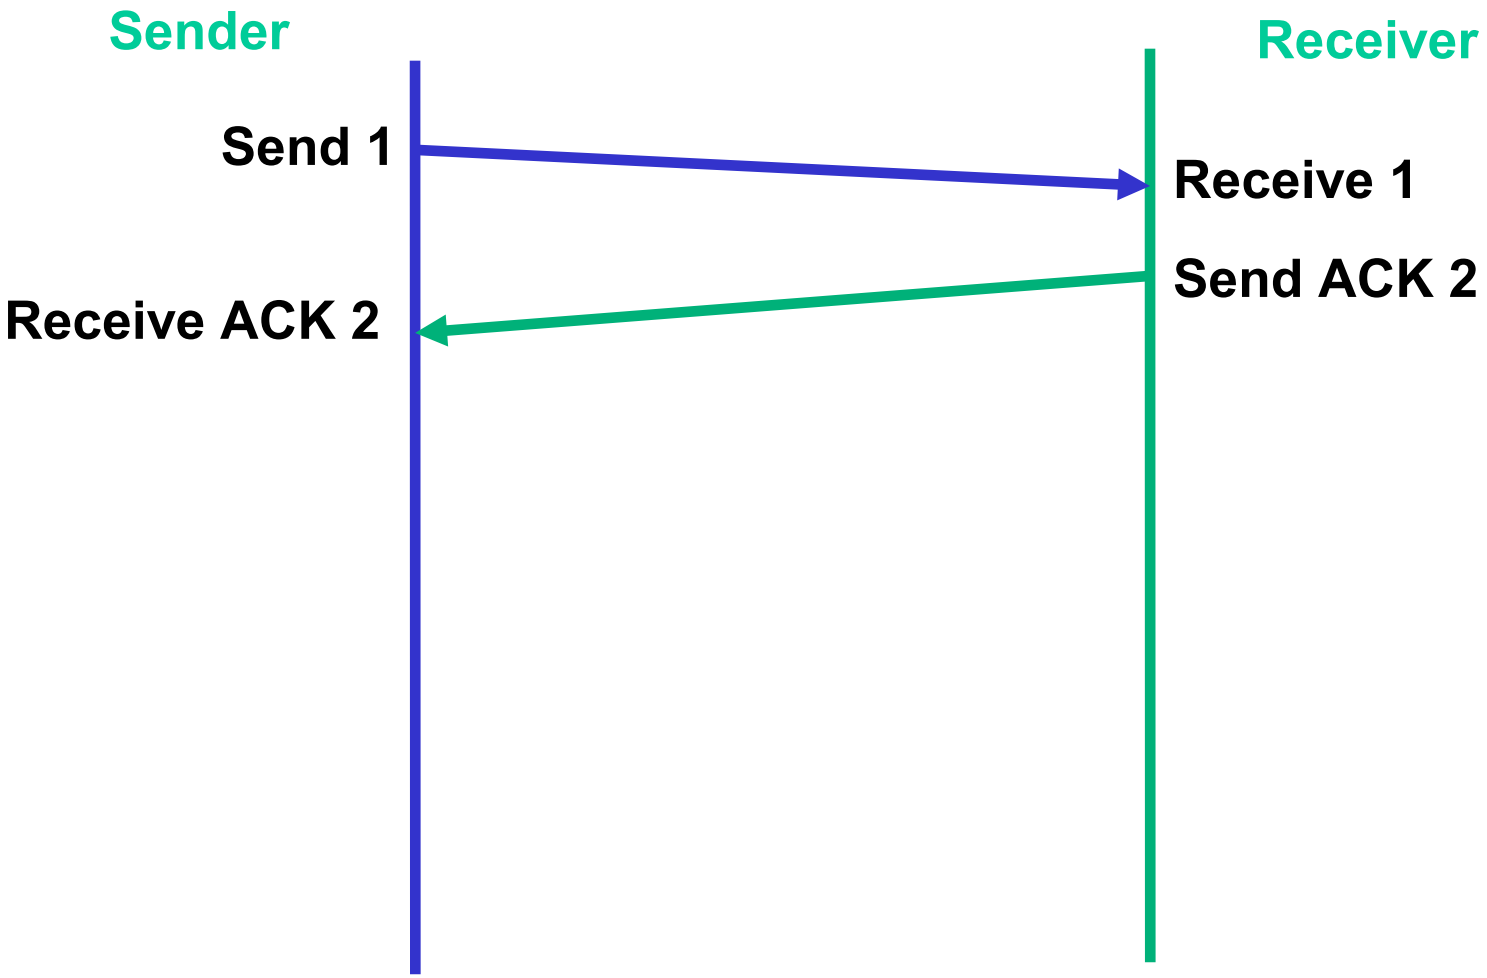

- Window size = 1

TCP Simple Acknowledgment

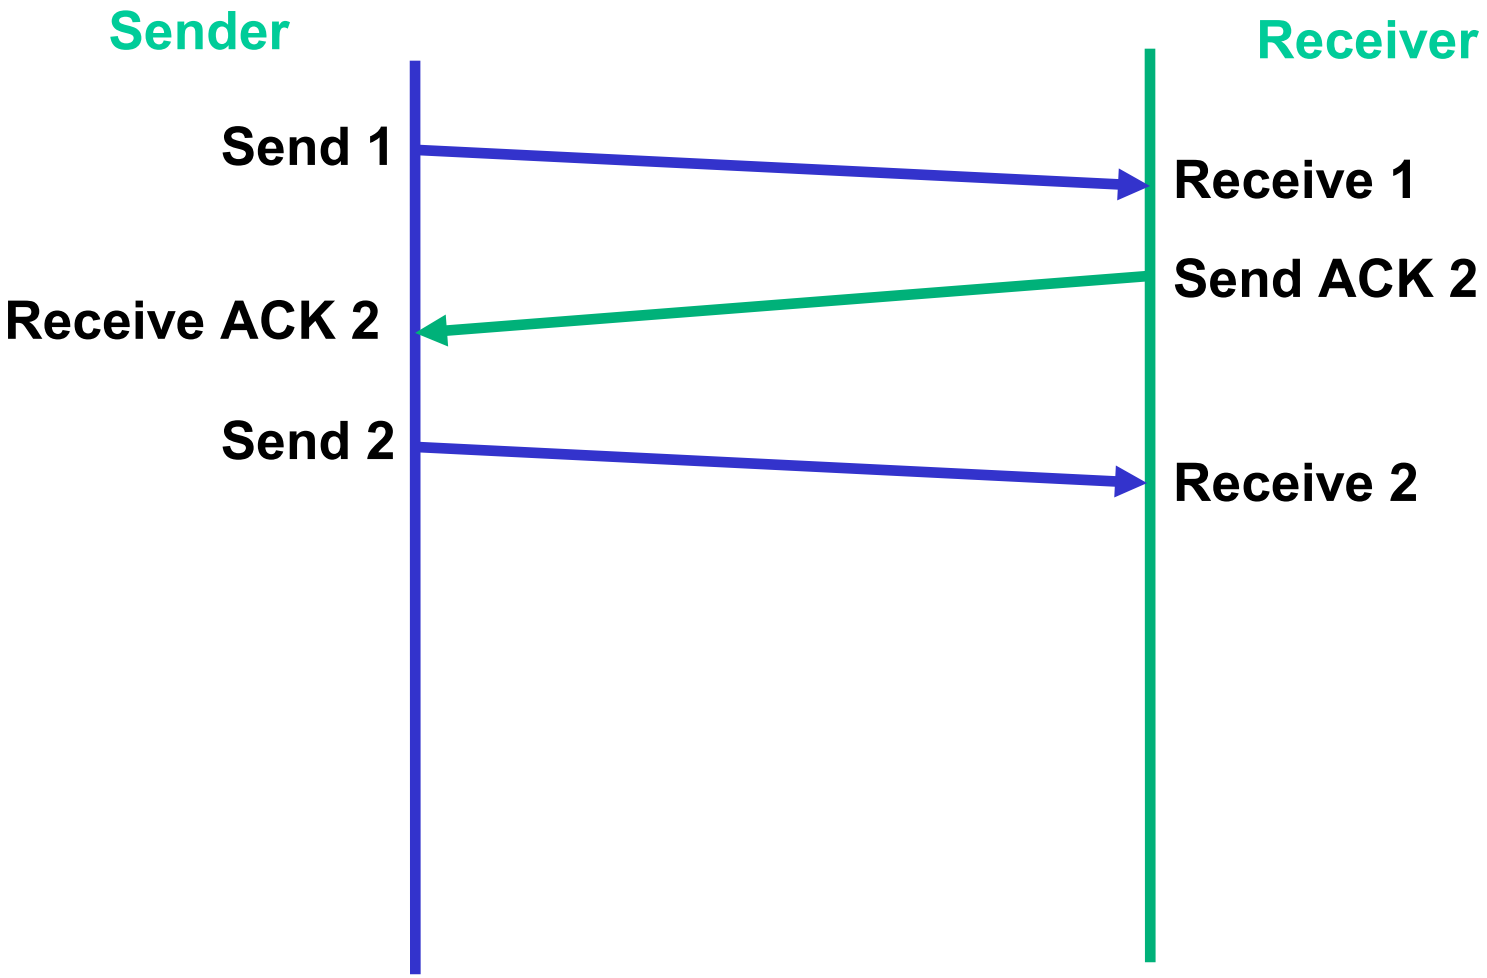

- Window size = 1

TCP Simple Acknowledgment

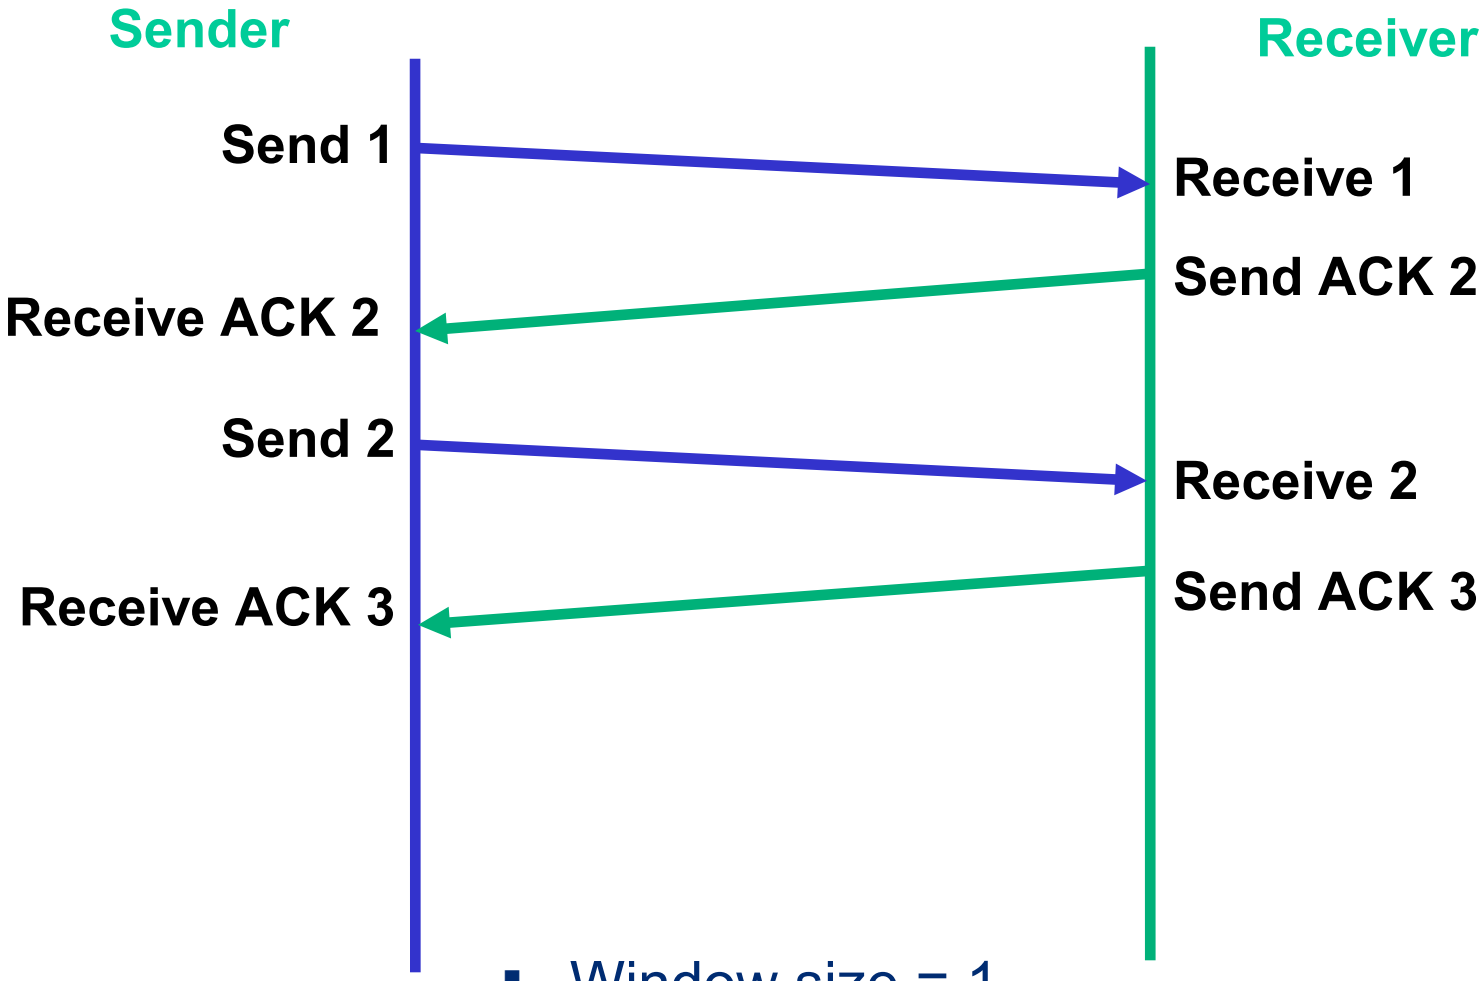

- Window size = 1

TCP Simple Acknowledgment

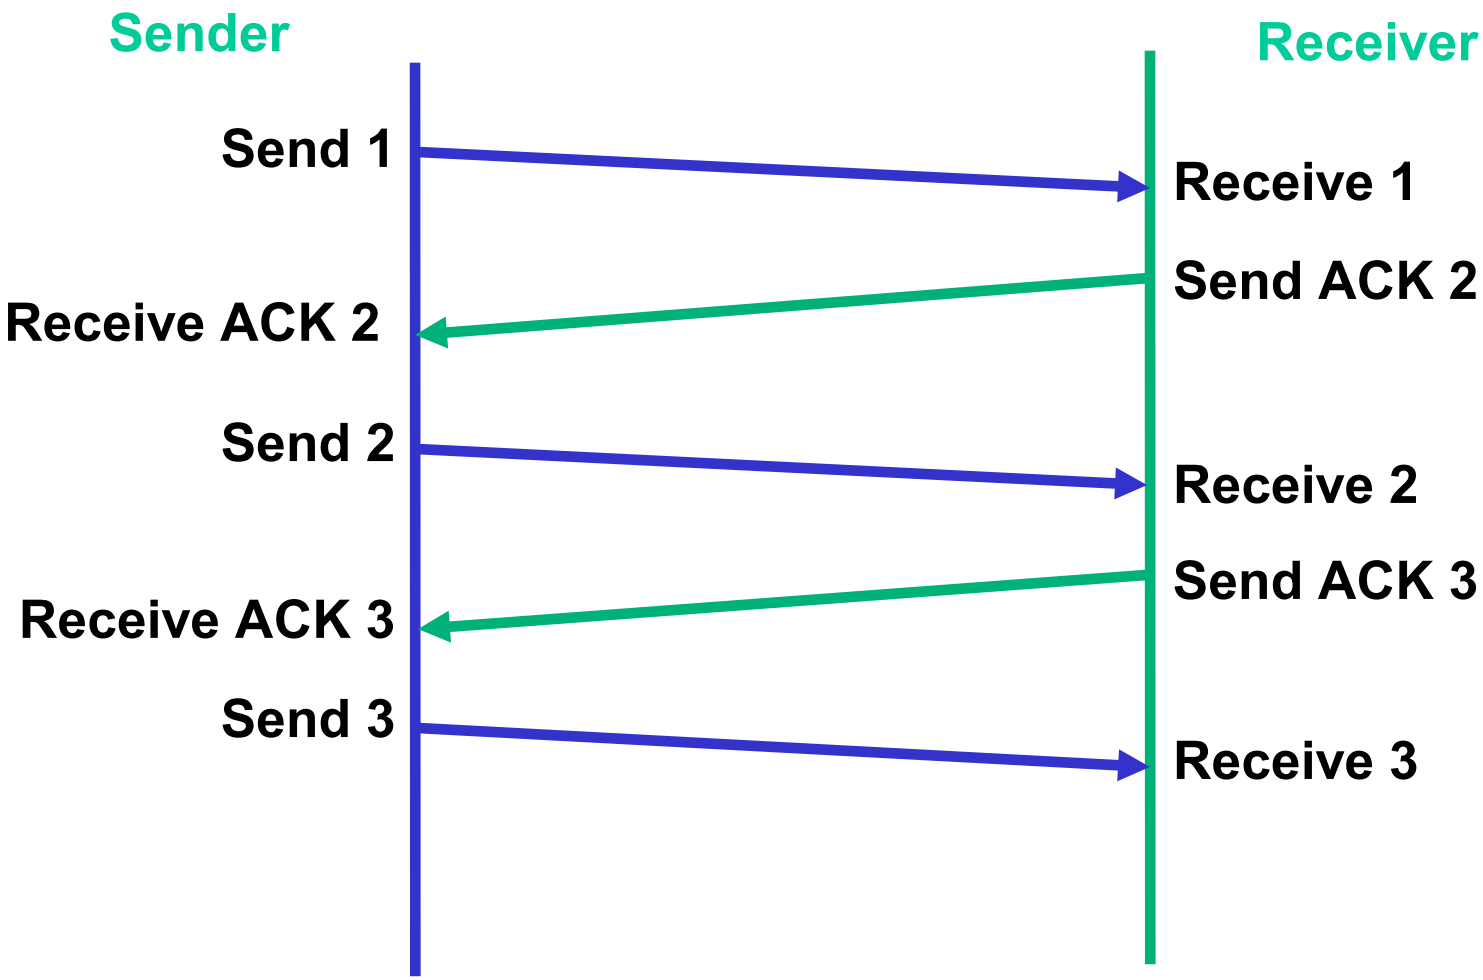

- Window size = 1

TCP Simple Acknowledgment

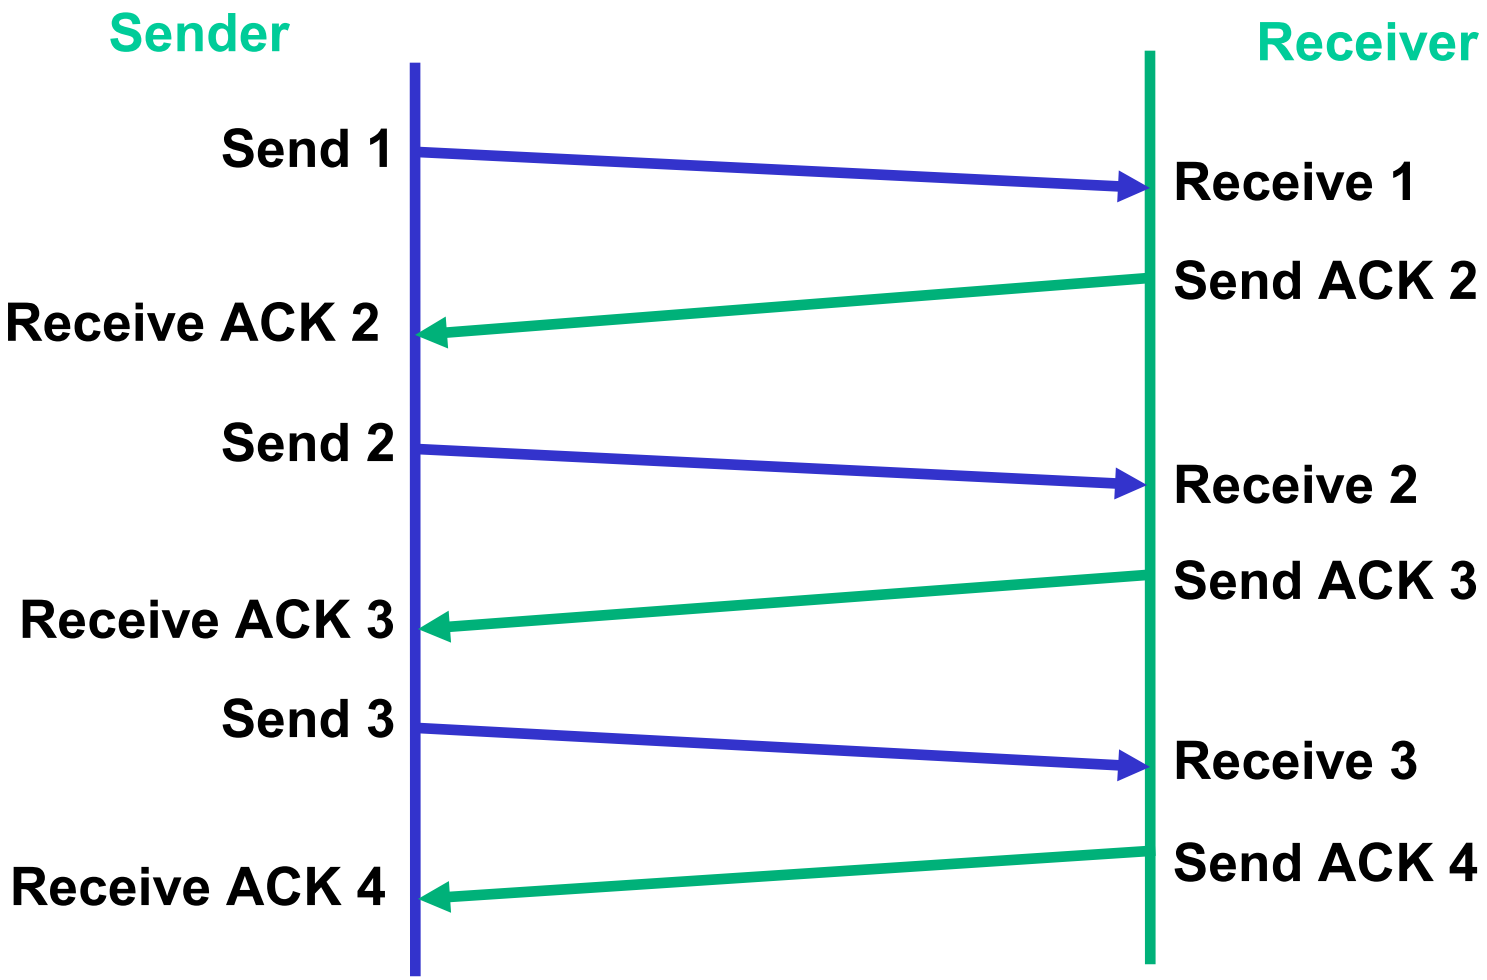

- Window size = 1

TCP Windowing

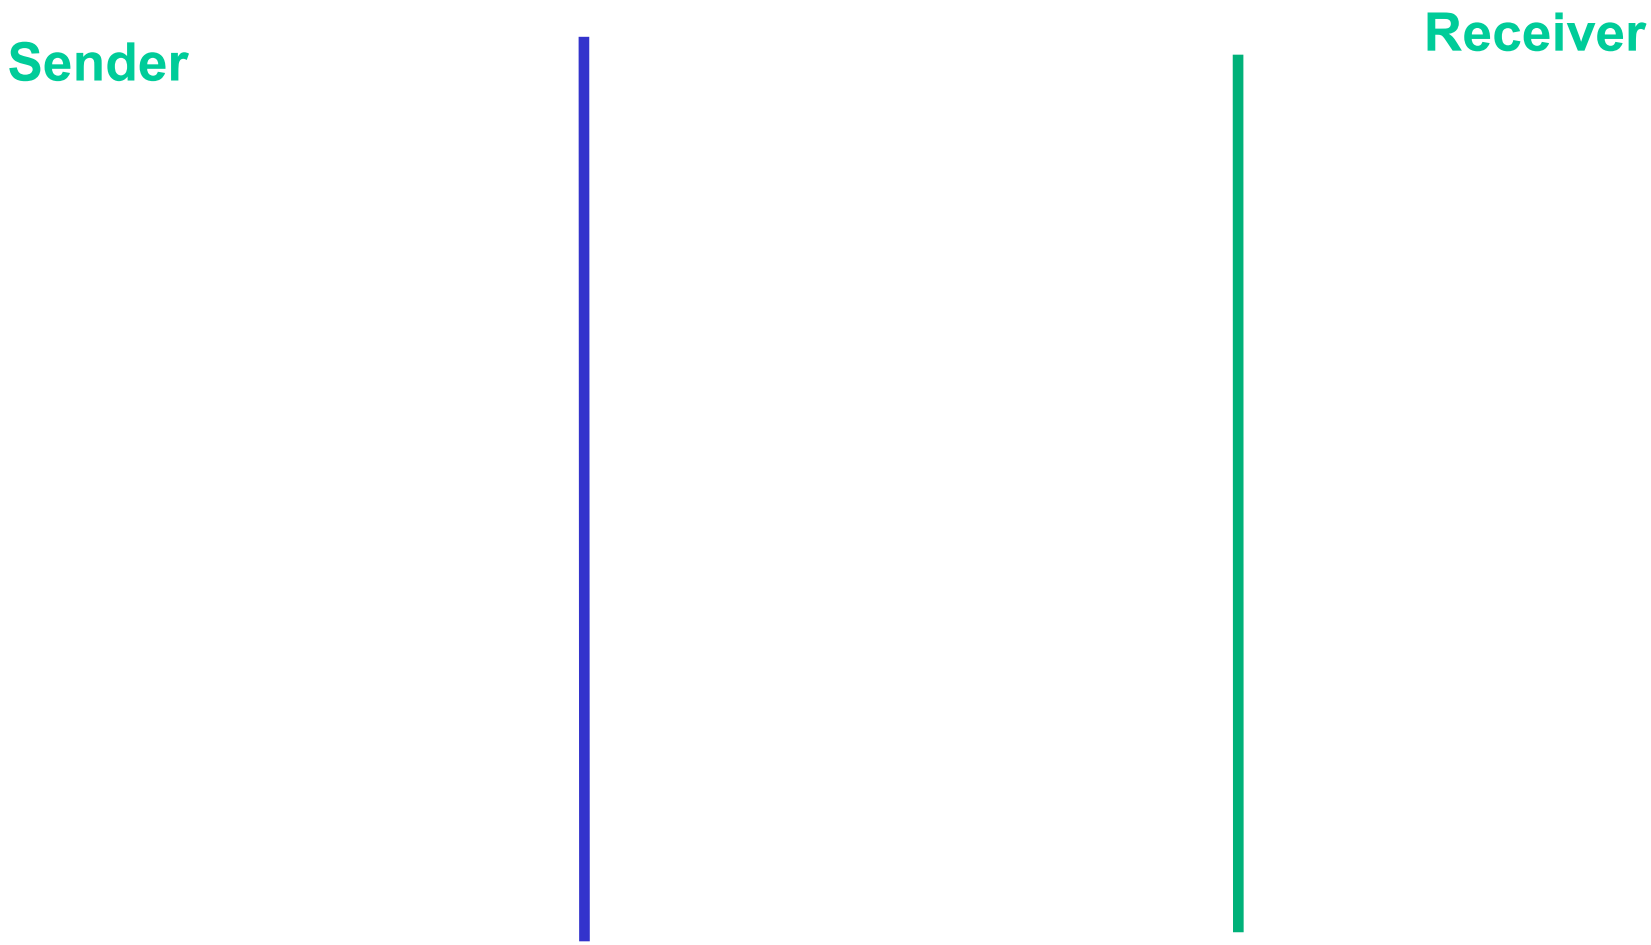

TCP Windowing

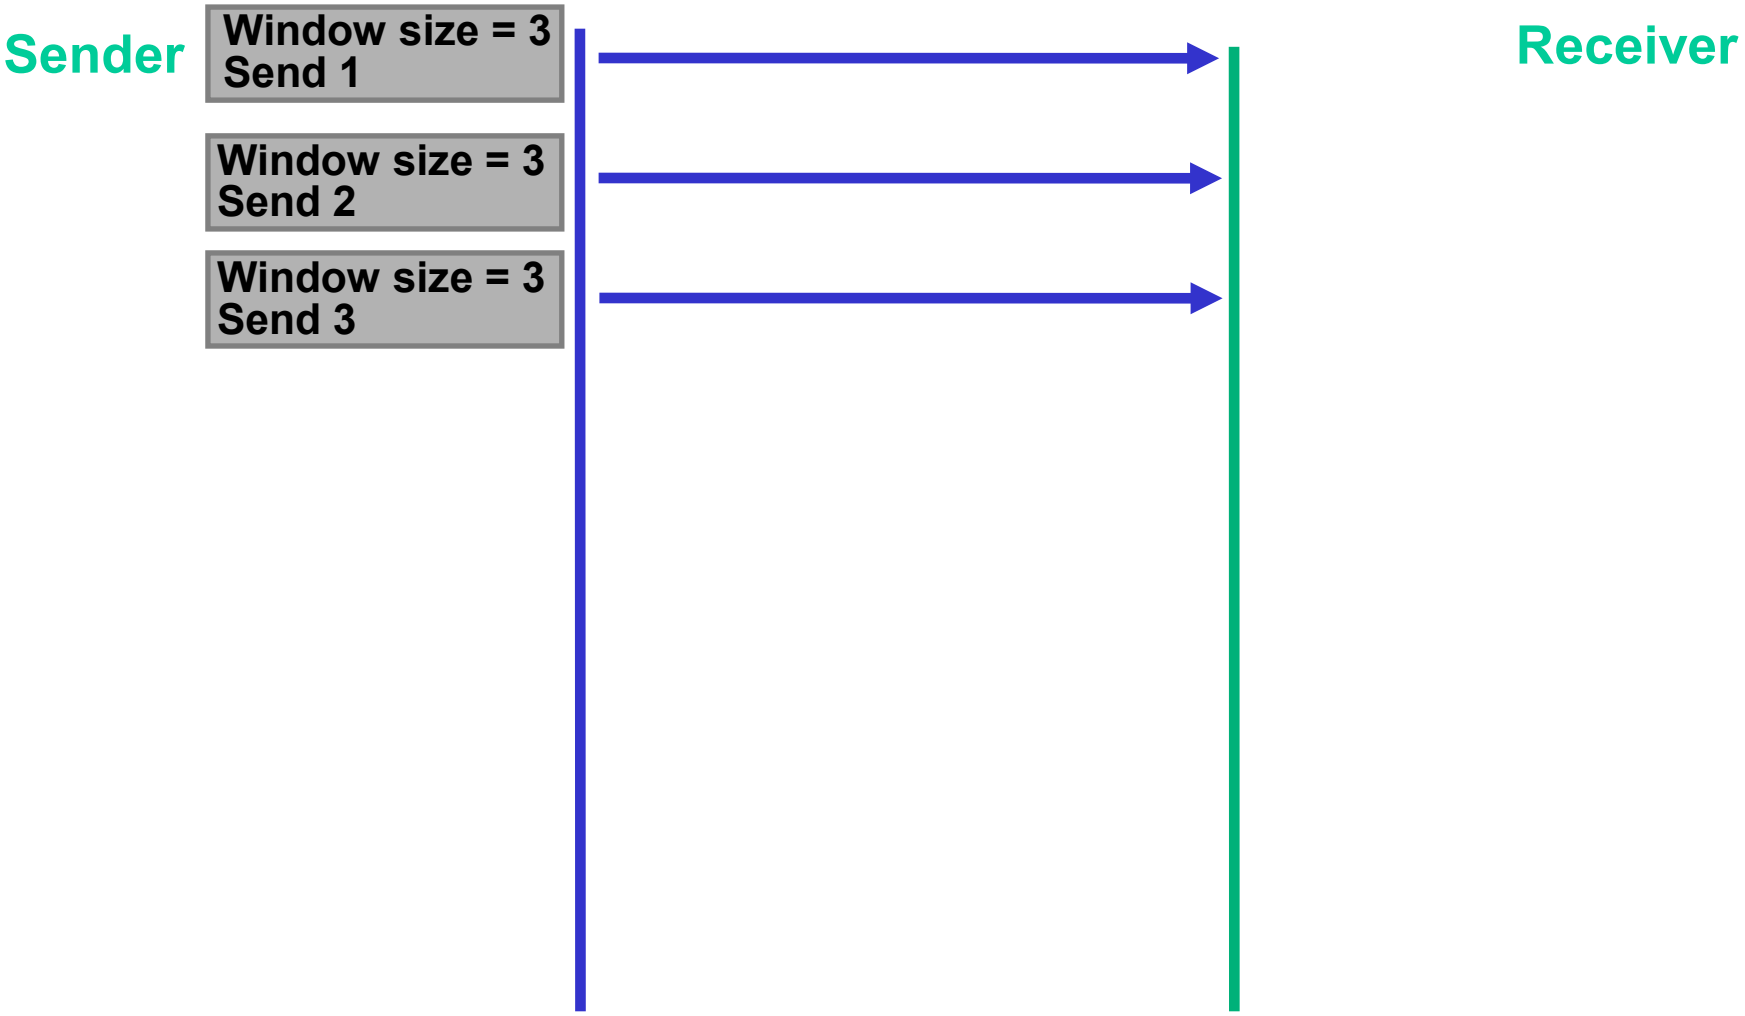

TCP Windowing

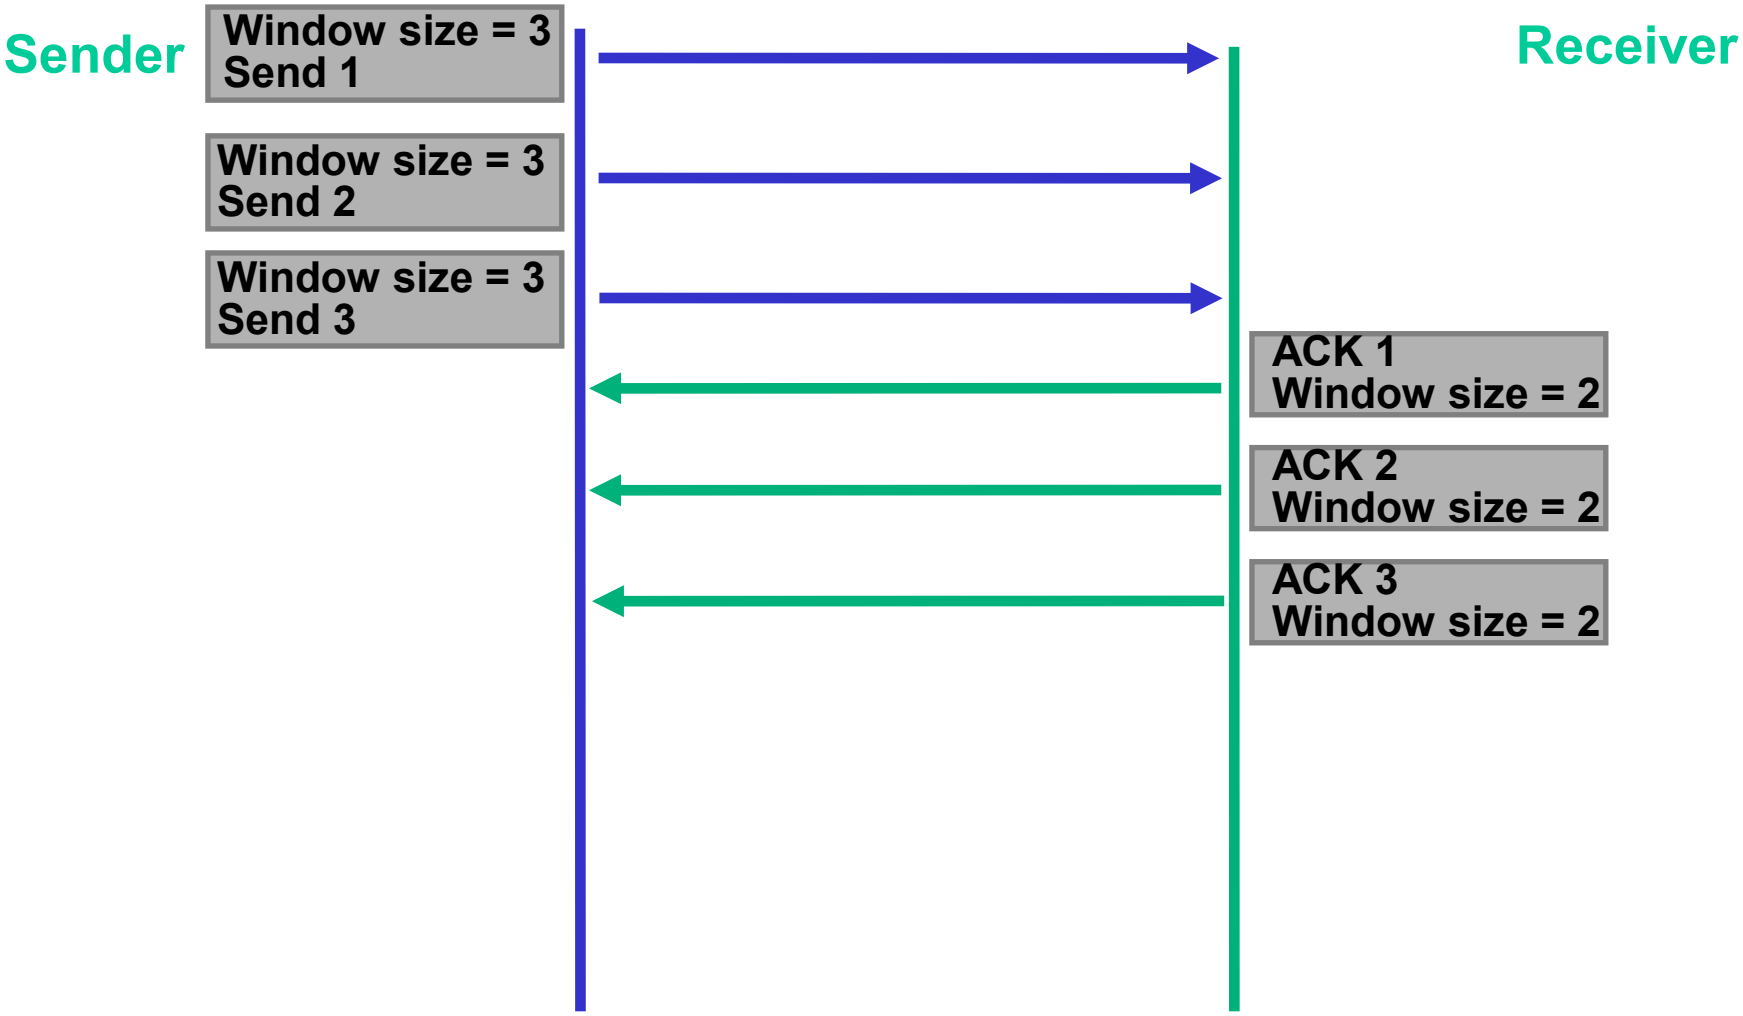

TCP Windowing

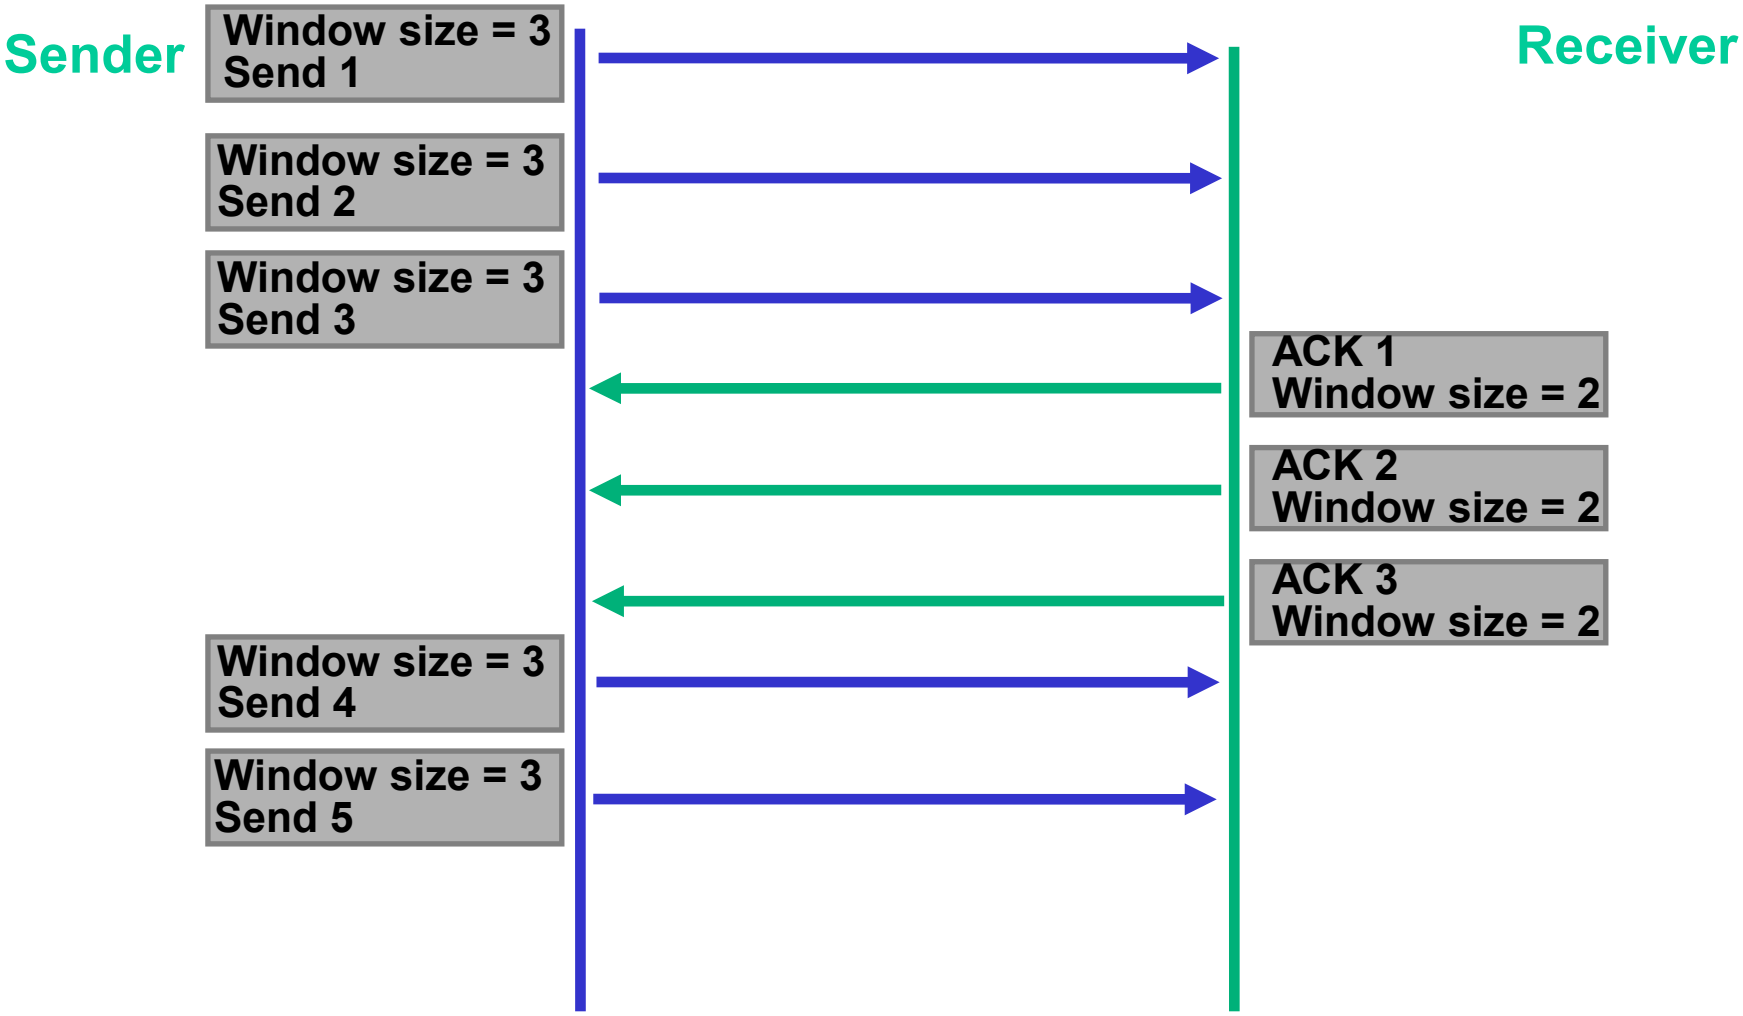

TCP Windowing

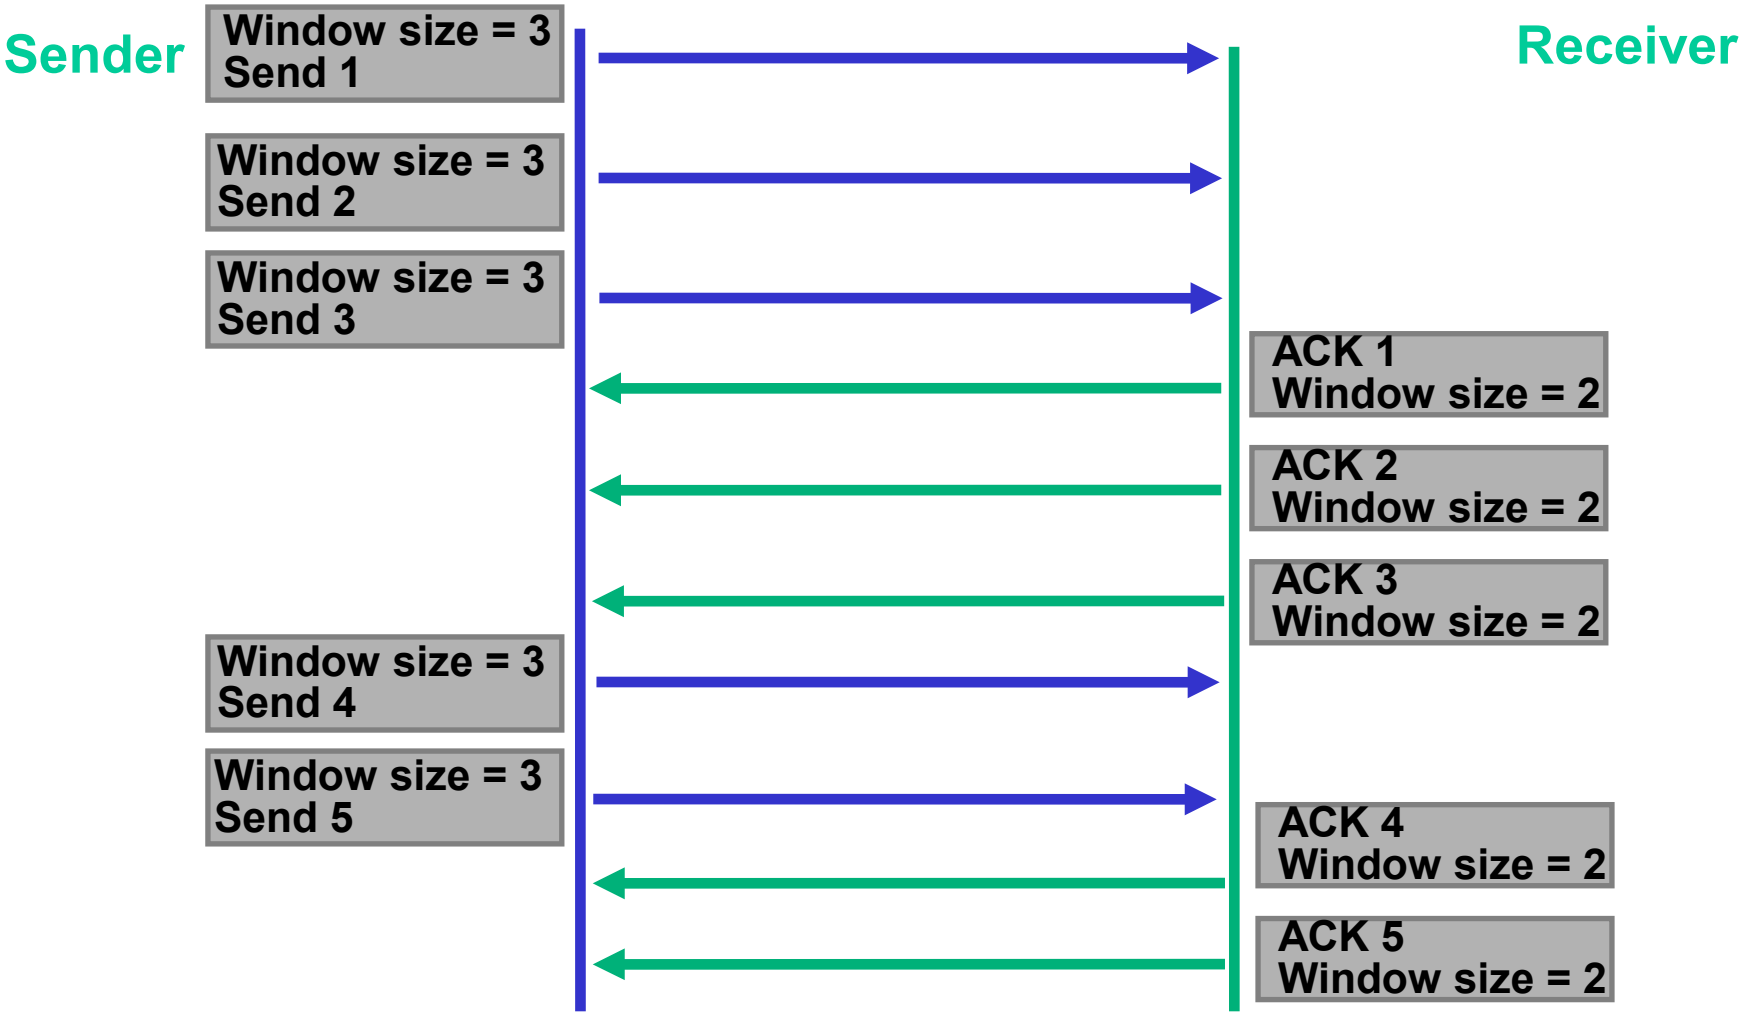

TCP Sequence and Ack Numbers

Source Port	Dest. Port	Sequence #	Acknowledgement #	…

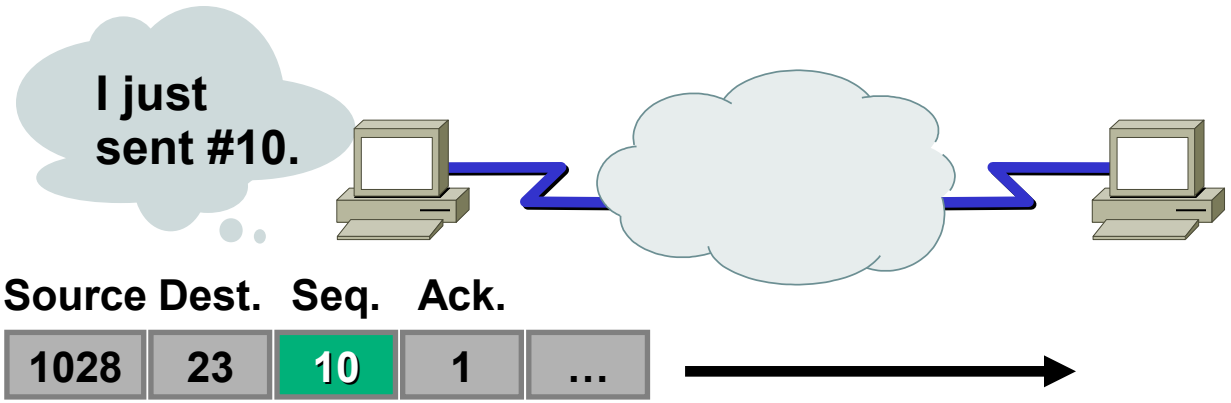

TCP Sequence and Ack Numbers

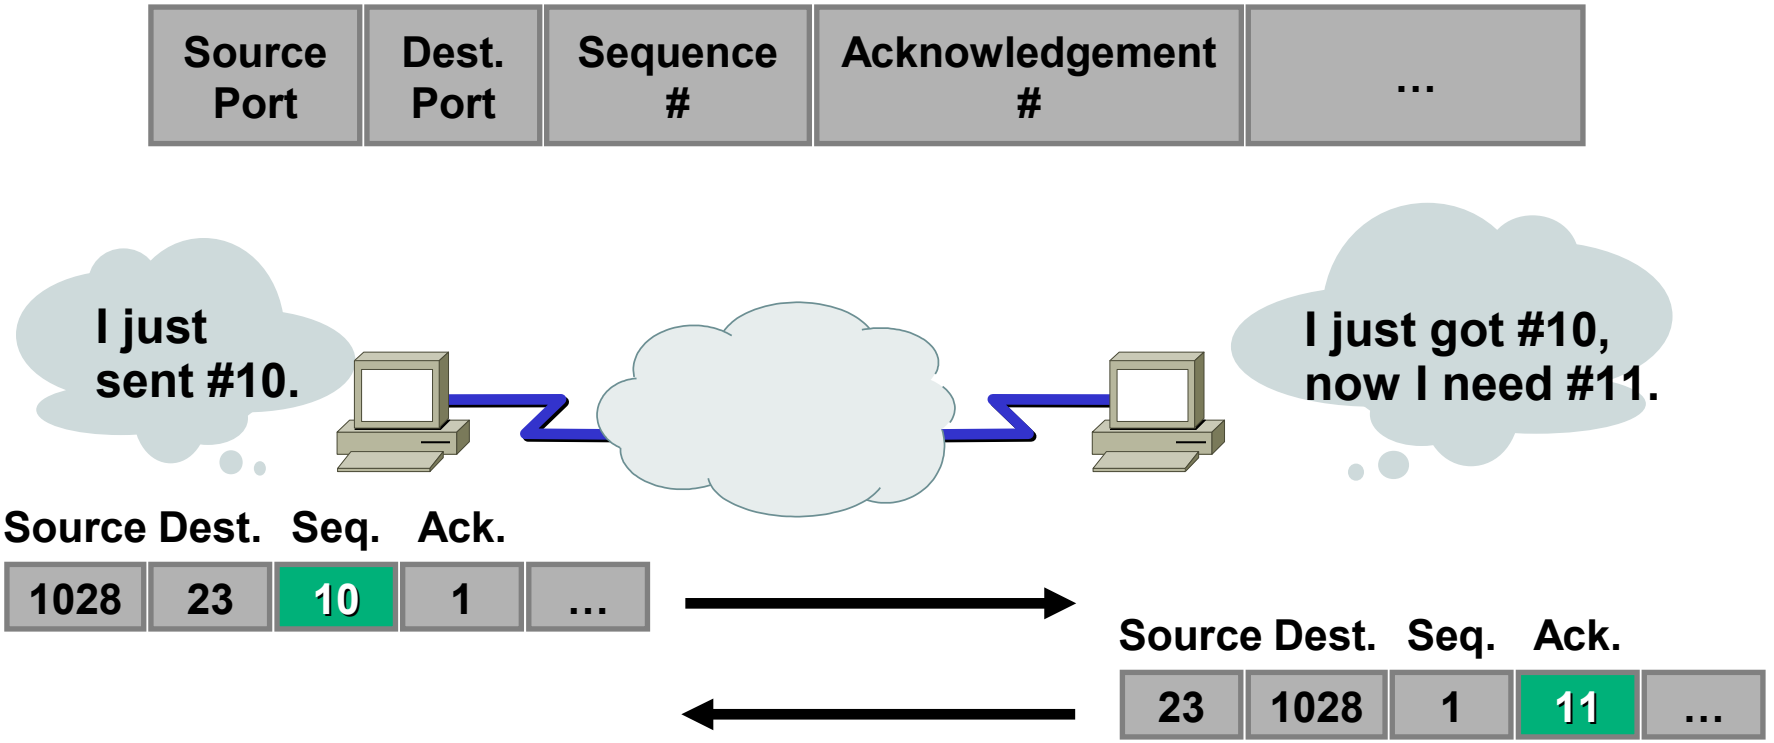

TCP Sequence and Ack Numbers

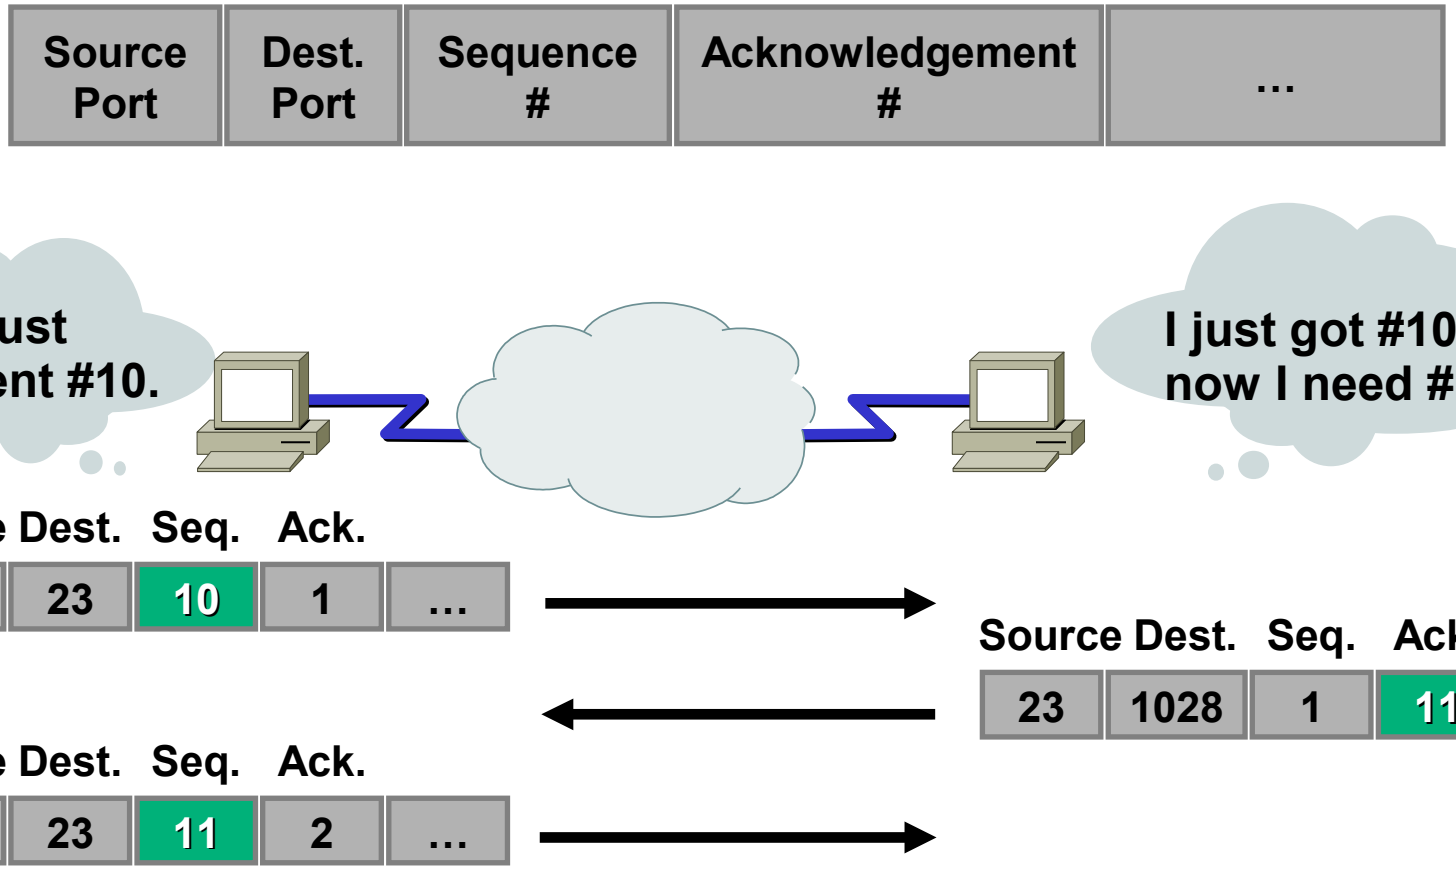

TCP Connection Closure

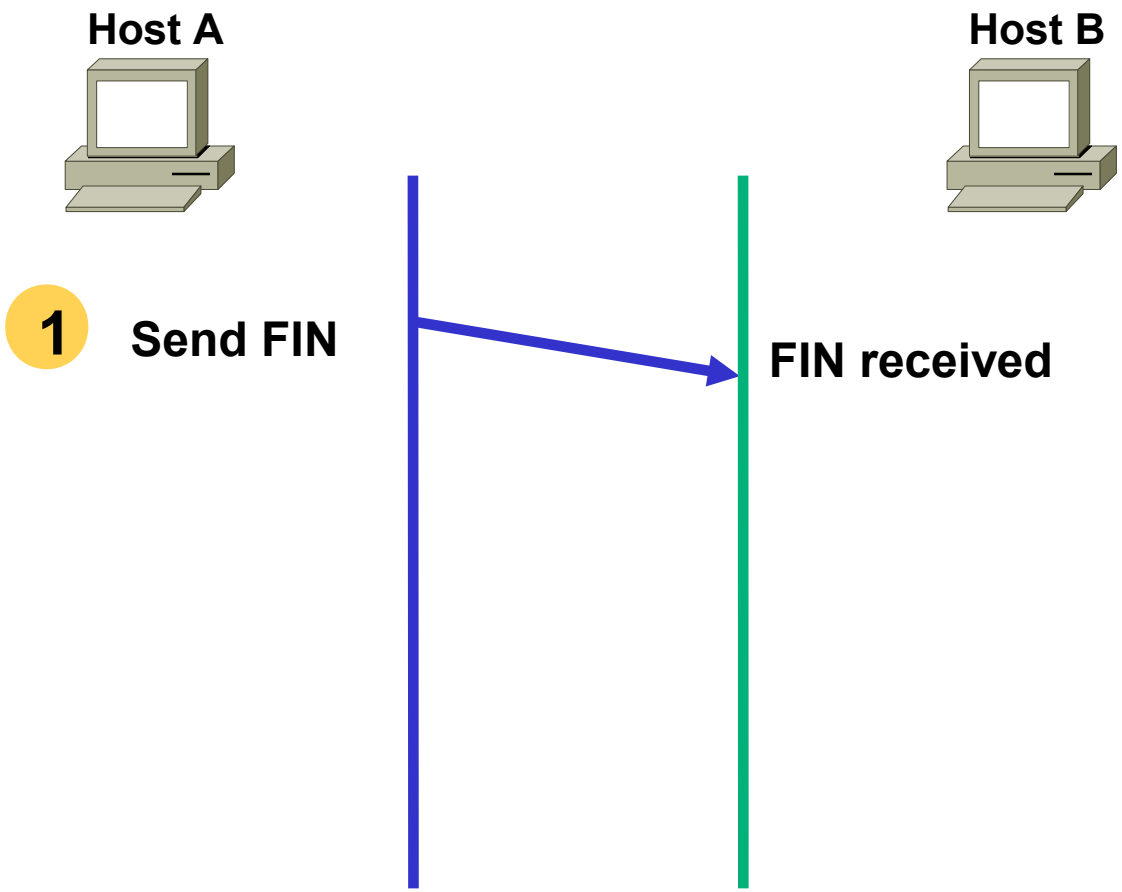

TCP Connection Closure

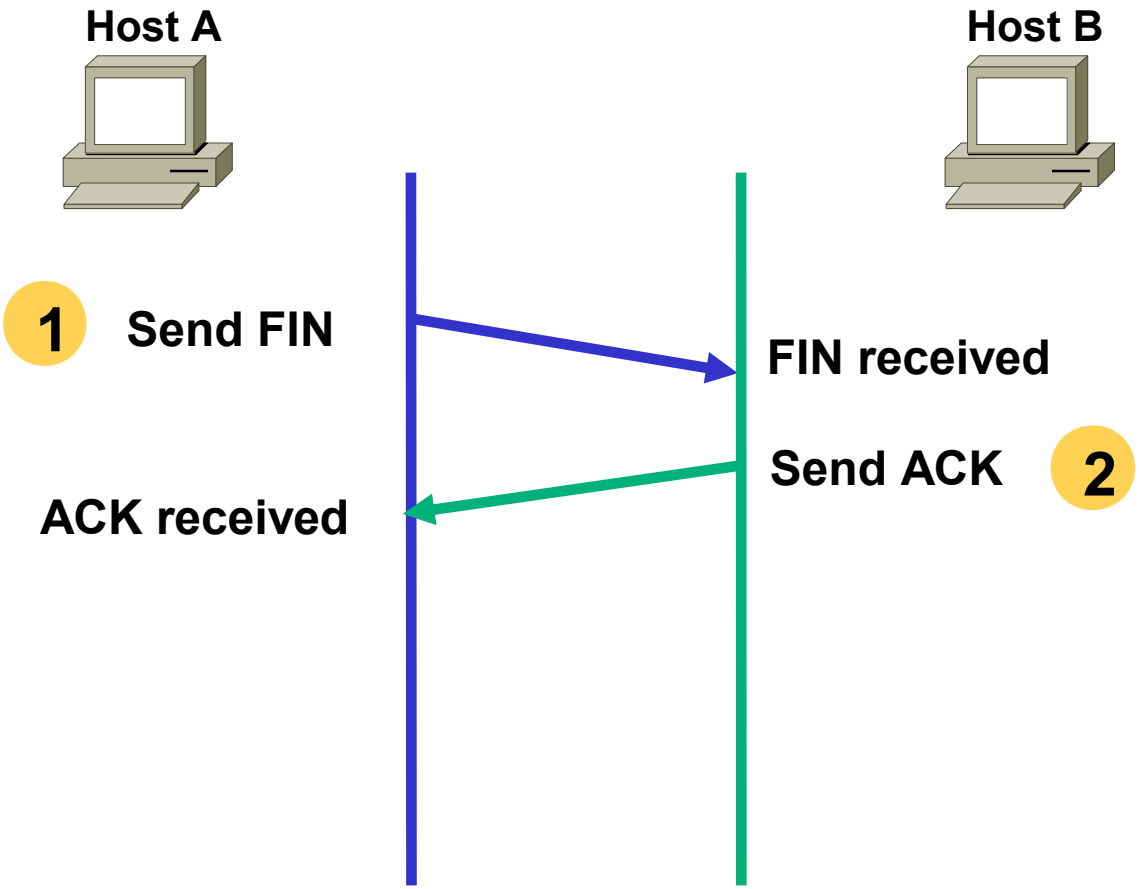

TCP Connection Closure

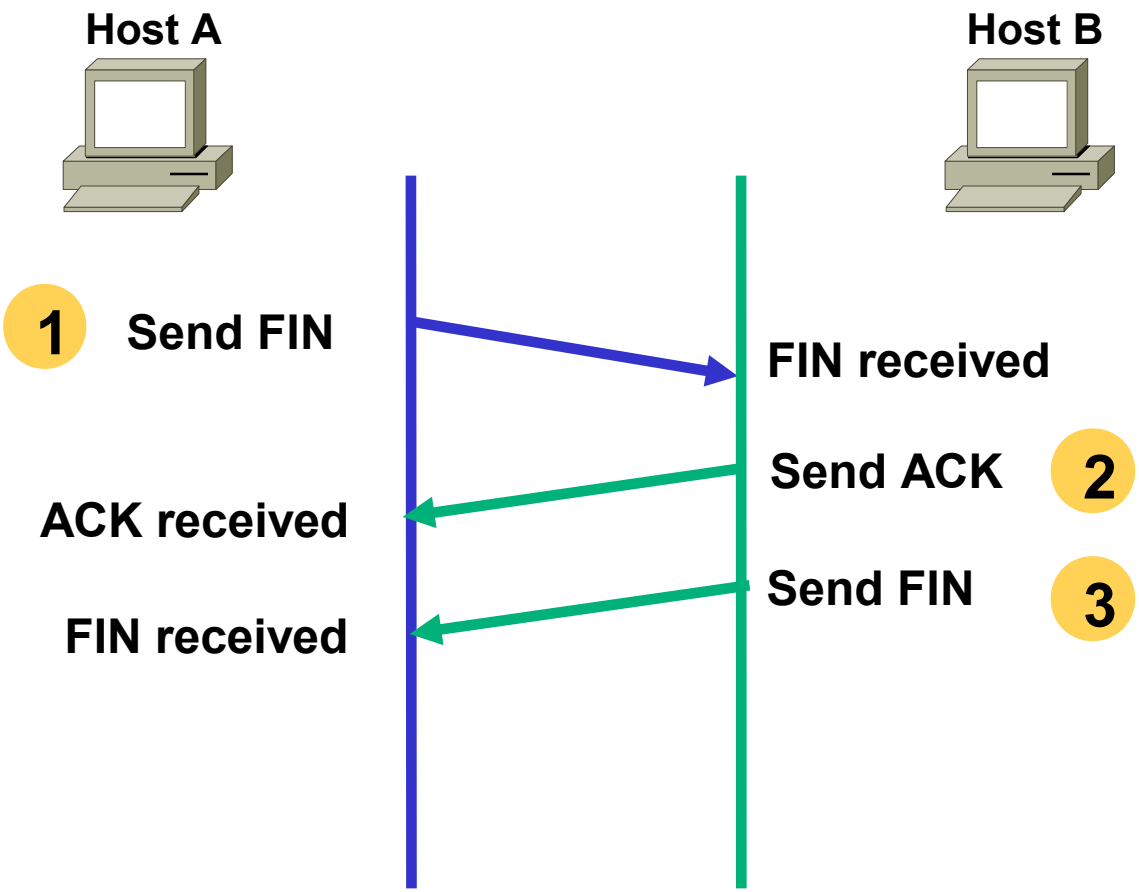

TCP Connection Closure

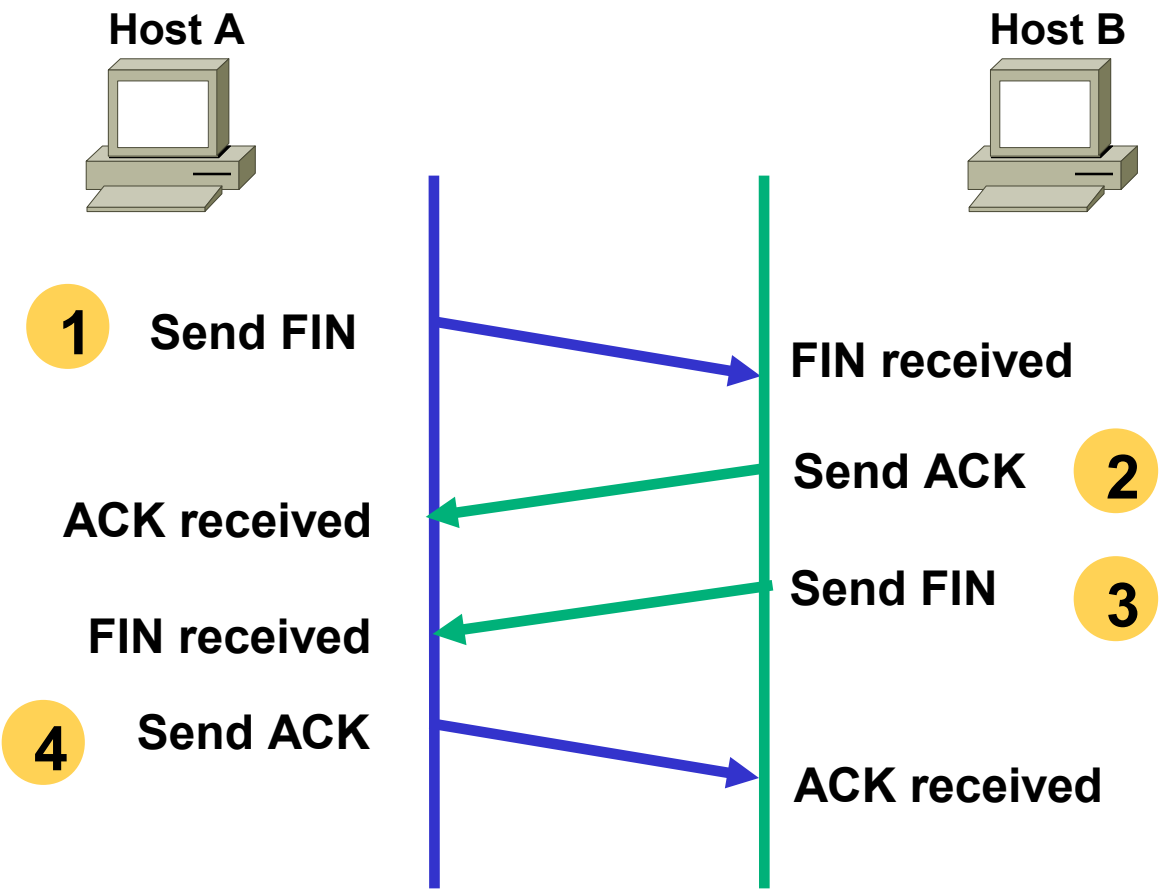

TCP Connection Closure

- Since TCP is full-duplex, connection must be shut down from both sides independently
  - it takes 4 segments to close the connection completely
- *Active close* – initiation of first FIN request
- *Passive close* – initiation of second FIN request

Connection closure sequence

- Application closes the session
- TCP sends FIN to the server
- Server TCP sends ACK to client
- Server TCP informs application
- Server application closes the session
- Server TCP sends FIN
- Client TCP sends ACK to server

TCP Connection Closure

- ***Half-close*** is also possible
  - Only one side discontinue transmission
  - Most of the present applications do not use *half-close*

Half-Close

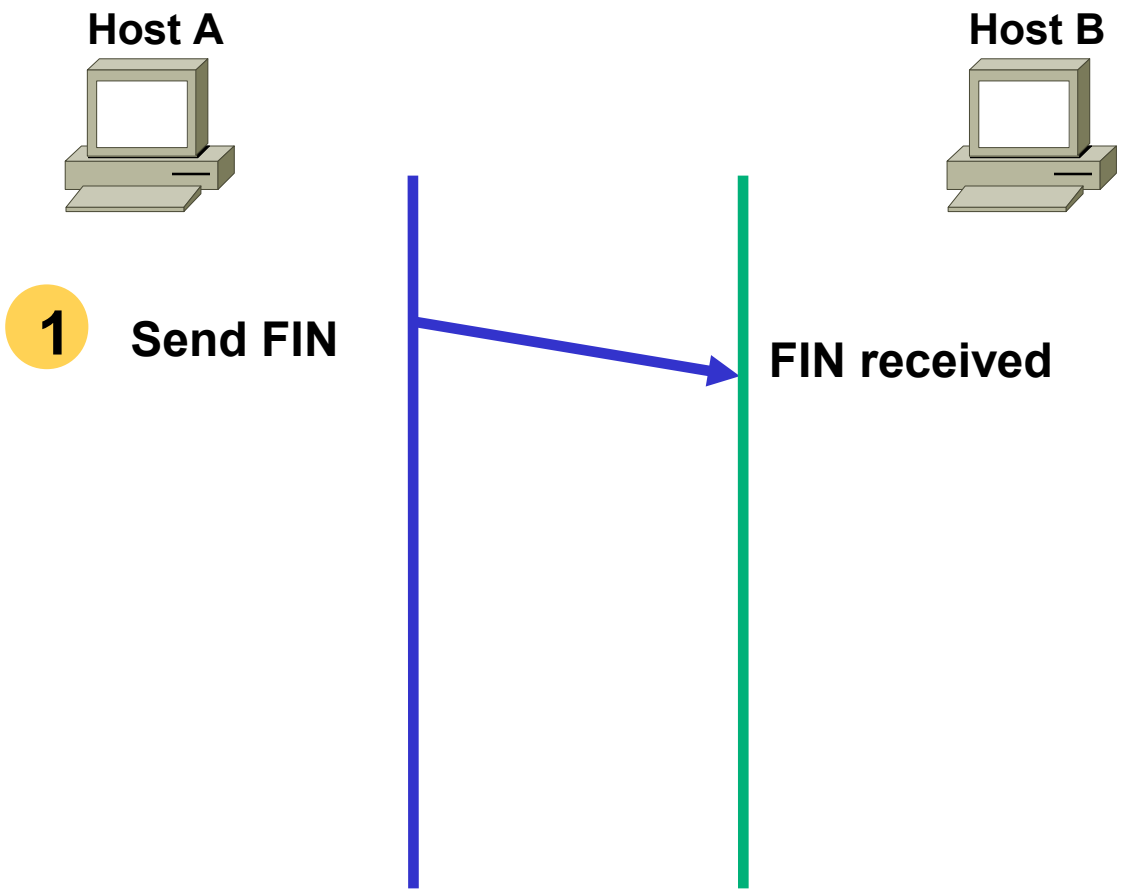

Half Close

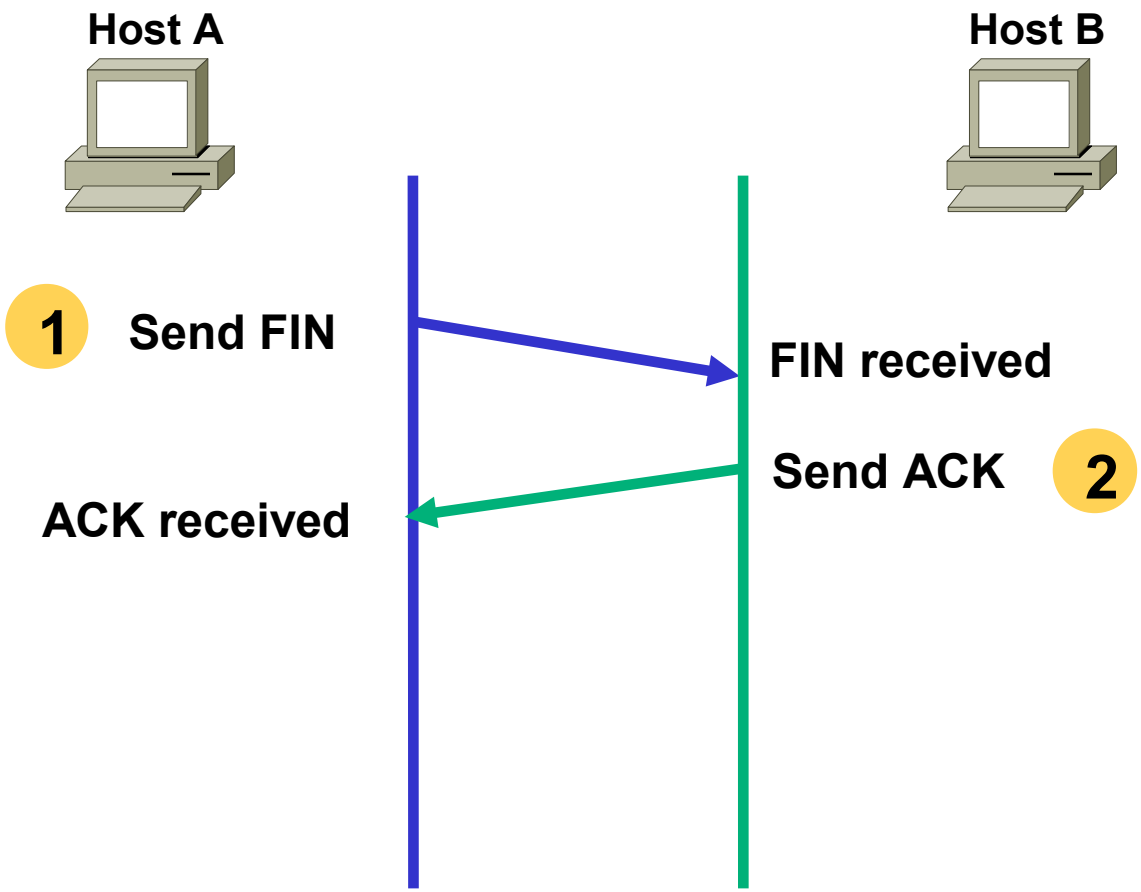

Half Close

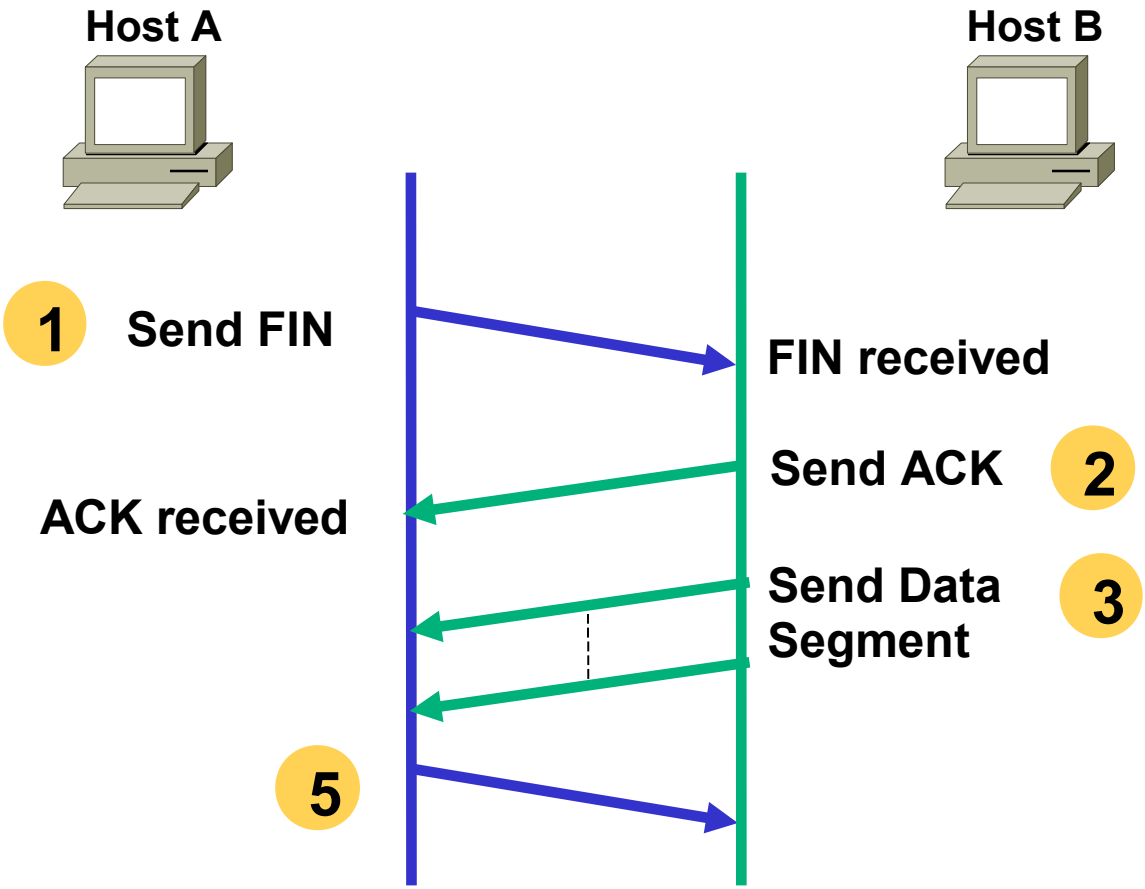
Host A
Host B
1
Send FIN
FIN received
Send ACK
2
ACK received
Send Data
Segment
3
5

Timeout during Connection Setup

- When initiator does not receive SYN-ACK
  - Resends SYN after some time
  - First retry after 6 sec
  - Second retry after 24 sec
  - Stops retrying after 75 seconds (Unix system)
  - These times are implementation dependent

MSS

- Maximum Segment Size
  - Exchanged with the initial SYN packets as an optional field
  - MSS does not appear in other packets
  - If MSS value is not received default value assumed is 536 bytes
  - 536+20IPhdr+20TCPHdr = 576byte IP datagram

MSS

- Larger the MSS, better it is for the protocol efficiency, provided it is not fragmented
- MSS may be at the most

  = MTU – 20 – 20 byte

  - For Ethernet = 1500-20-20 = 1460
- If destination is non-local, MSS normally defaults to 536
- MSS is configurable value

MSS

- When two sides announce different MSS, they normally settle down to the lower value.
  - This is not mandatory
  - Avoids fragmentations
  - Not necessarily eliminate fragmentations if intermediate links have even lower MTU
  - Use MTU discovery mechanism to avoid this

Non-delivery of TCP Segments

- Non delivery is indicated by non-receipt of ACK at sender side
- This may be caused by
  - Loss of packet
    - ❖ Due to congestion
    - ❖ Due to error in header
    - ❖ Due to error in data
  - Loss of ACK
  - Delayed delivery by IP

Non-delivery of TCP Segments

- Non-receipt of ACK is decided by
  - Retransmission timer
  - Receipt of duplicate ack
- TCP assumes that the non-delivery is because of congestion
- Reduces the window size when the packet is declared as undelivered
  - Not an efficient way if loss of packet is due to a transmission error

Types of Connection Closure

- *Orderly release* – graceful shutdown
  - When closure initiated by applications
  - No loss of data
  - Using *FIN* segment
- *Abortive release*
  - Abrupt termination
  - Using reset (*RST*)

Variation in Connection Open and Close

- Simultaneous Opening
  - Both sides send SYN
  - Both sides respond with SYN-ACK
- Simultaneous Closure
  - Both sides send FIN
  - Both sides send ACK

RST

- Generated on receipt of an ***incorrect*** TCP segment
  - Packet does not belong to the referenced connection, determined by
    - IP
    - Port number
    - Sequence number
- Generated on receipt of connection request to an nonexistent port
- Generated by application when it aborts the application

RST

- At sender side, any queued data is thrown away
- At the receiver side, APIs used should be able to inform application about the abortive release

PSH

- PUSH Flag
- Indicates to the receiver to send the data to the application without further delay
- Used in the interactive applications or during interactive operations
- Also used when last portion of the data is sent by sender stack

Half-open Connections

- One side abruptly terminates the session
- May be caused by
  - System crash
  - Machine powered off without graceful shutdown
- Server will not know the closure and will be in wait state
- Security risk

Interactive Data Flow

- ***Tinygrams***
  - Small data flow during interactive applications
  - Example: ***Rlogin***
- Nagle Algorithm (RFC 896)
  - TCP connection can have only one outstanding small segment that is not yet acknowledged
  - Small data is collected by TCP and sent together when *Ack* is received for previous small segment
  - Might cause problem when ASCII escape character is involved (special function keys)

Interactive Data Flow

- Repacketisation
  - Sending retransmission and next segment in the same segment
- TCP protocol spoofing
  - Used when delay is high, bandwidth is sufficent
  - Improves user experience and application performance
  - Used in VSATs

Sliding Window

- Start small – ***slow start***
- Grow exponentially
- Bound by upper limit of window size
- Reduces window size when encounter segment loss
- Increases window size again
  - Offered window size
  - Usable window size

Sliding Window

- *Slow Start*
  - *Rate of transmission depends on rate of receipt of acknowledgments*
  - *A flow control imposed by sender based on its assessment of congestion in the network*

Congestion Avoidance

- Indication of loss of packet
  - Timeout
  - Duplicate ACK
- On receiving duplicate ACK reduce usable window size to half
- If congestion is indicated by timeout, reduce usable window size to one, initiate slow start

Fast Retransmit

- When three or more duplicate ACK received, retransmit the un-ACKed packets without waiting for timeout of retransmission timer

Reading Assignment

- Congestion Avoidance
- Slow Start
- Fast Retransmit
- Fast Recovery

Congestion Avoidance

- Indication of loss of packet
  - Timeout
  - Duplicate ACK
- Two variables
  - ***cwnd*** (congestion window)
  - ***ssthresh***(slow start threshold)
- On congestion, ***ssthresh*** = ***cwnd***/2
- If congestion is indicated by timeout, ***cwnd*** is set to one, slow start

Thank you!